

Big Data Analysis Platforms

SHYI-CHYI CHENG

Slides credited to Cloudera Developer Training for Apache Hadoop

Outline

- Review of Virtual Machine (虛擬機器回顧)
- Hadoop Platform (運算分析系統架構)
- **MapReduce**
- Introduction to Python (Python入門簡介)
- Python Spark Platform (Python Spark運算分析架構)
- Parallel Programming With Spark

MapReduce

Original Slides by
Owen O'Malley (Yahoo!)
&
Cloudera Developer Training for Apache Hadoop

Introduction to MapReduce

- What you will learn
 - The concepts behind MapReduce
 - How data flows through MapReduce stages
 - Typical uses of Mappers
 - Typical uses of Reducers

MapReduce - What?

- MapReduce is a programming model for efficient distributed computing
- It works like a Unix pipeline
 - `cat input | grep | sort | uniq -c | cat > output`
 - **Input** | **Map** | Shuffle & Sort | **Reduce** | **Output**
- Efficiency from
 - Streaming through data, reducing seeks
 - Pipelining
- A good fit for a lot of applications
 - Log processing
 - Web index building

MapReduce Features (1)

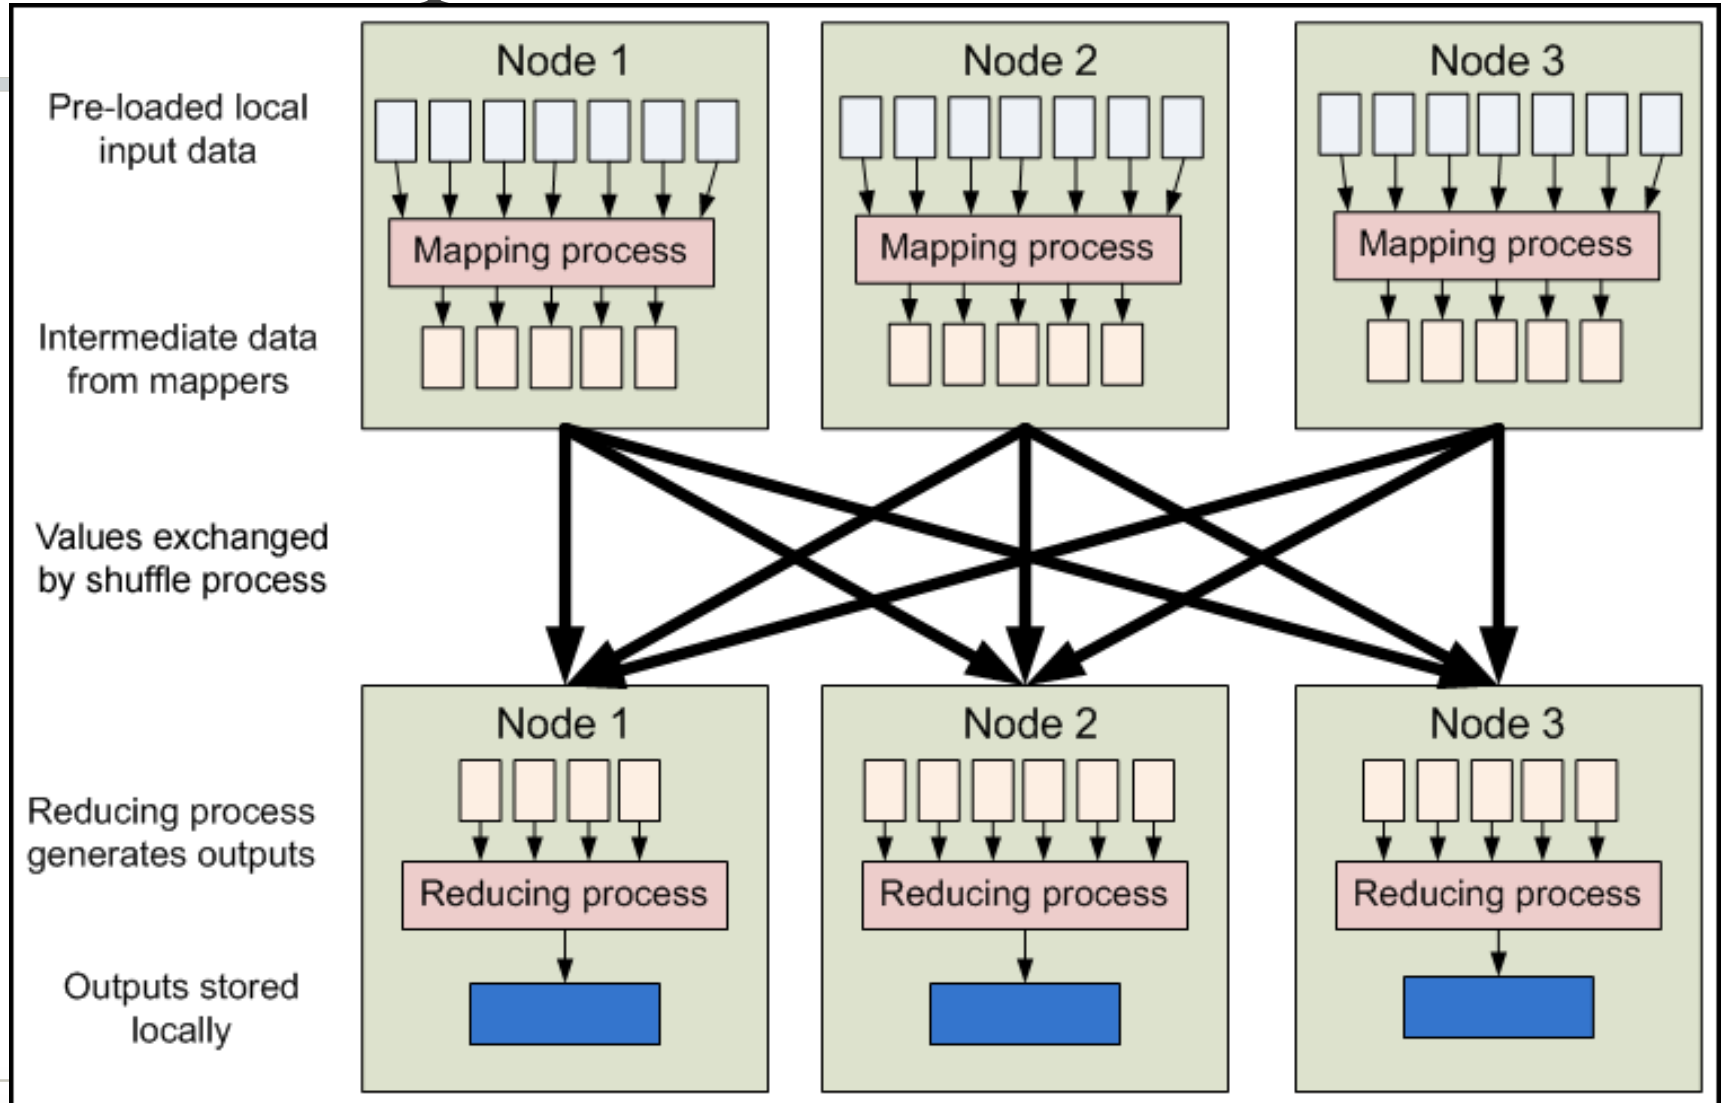
- Automatic parallelization and distribution
- Fault-tolerance
- A clean abstraction for programmers
 - MapReduce programs are usually written in Java
 - Can be written in any language using **Hadoop streaming**
 - All of Hadoop is written in Java
- MapReduce abstracts all the “housekeeping” away from the developer
 - Developer can simply concentrate on writing the Map and Reduce functions

Hadoop Streaming is a utility which allows users to create and run jobs with any executables (e.g. shell utilities) as the mapper and/or the reducer.

MapReduce Features (2)

- Fine grained Map and Reduce tasks
 - Improved load balancing
 - Faster recovery from failed tasks
- Automatic re-execution on failure
 - In a large cluster, some nodes are always slow or flaky
 - Framework re-executes failed tasks
- Locality optimizations
 - With large data, bandwidth to data is a problem
 - Map-Reduce + HDFS is a very effective solution
 - Map-Reduce queries HDFS for locations of input data
 - Map tasks are scheduled close to the inputs when possible

MapReduce - Dataflow



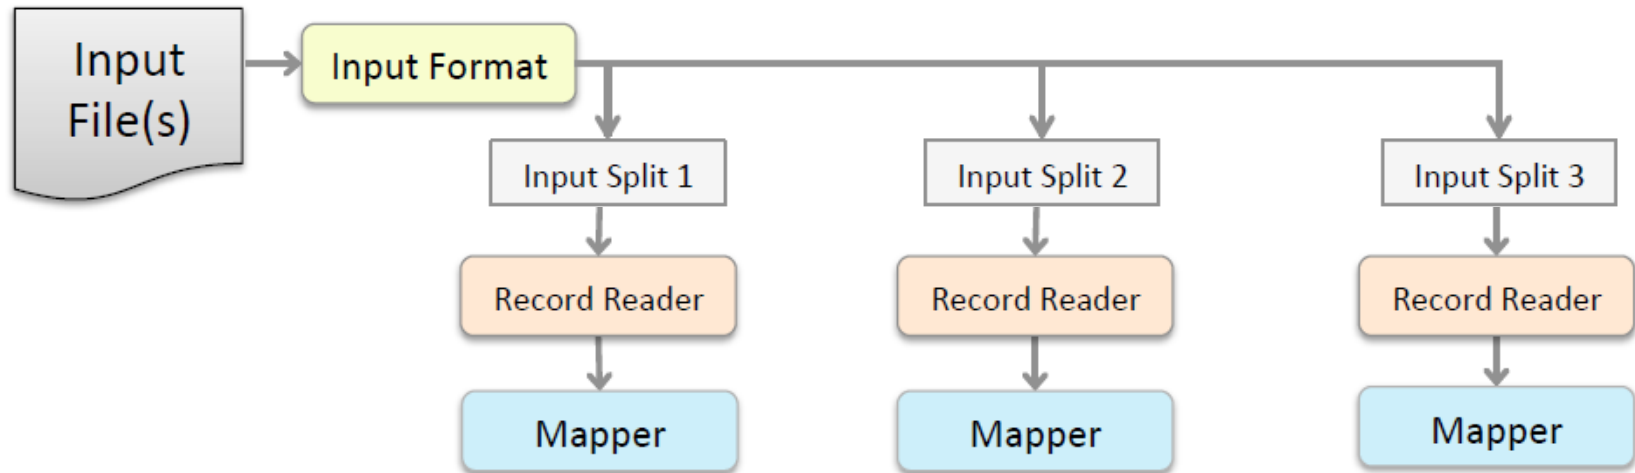
Key MapReduce Stages

- The Mapper
 - Each Map task (typically) operates on a single HDFS block
 - Map tasks (usually) run on the node where the block is stored
- Shuffle and Sort
 - Sorts and consolidates intermediate data from all mappers
 - Happens after all Map tasks are complete and before Reduce tasks start
- The Reducer
 - Operates on shuffled/sorted intermediate data (Map task output)
 - Produces the final output

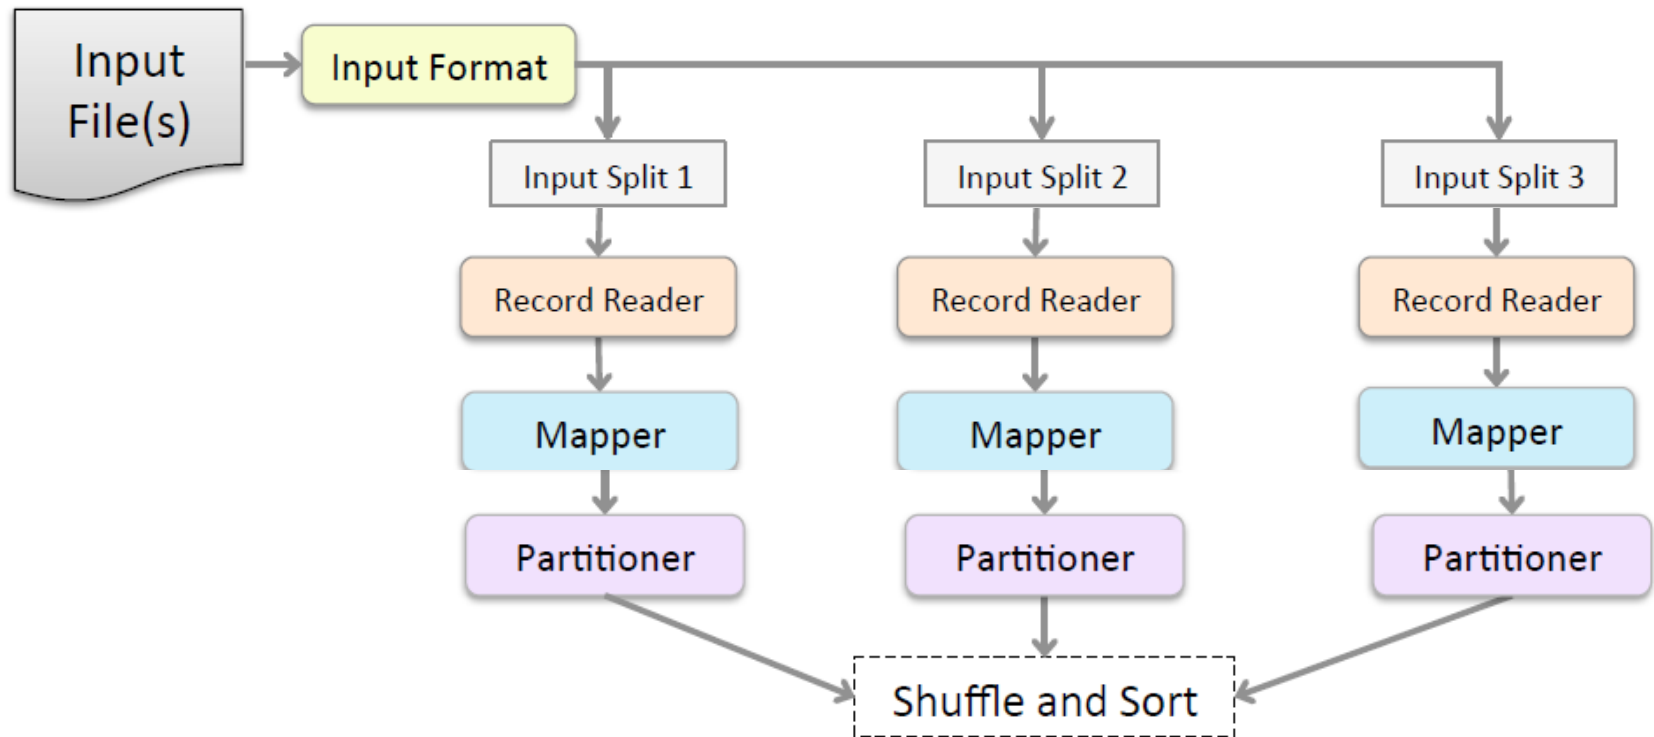
The MapReduce Flow



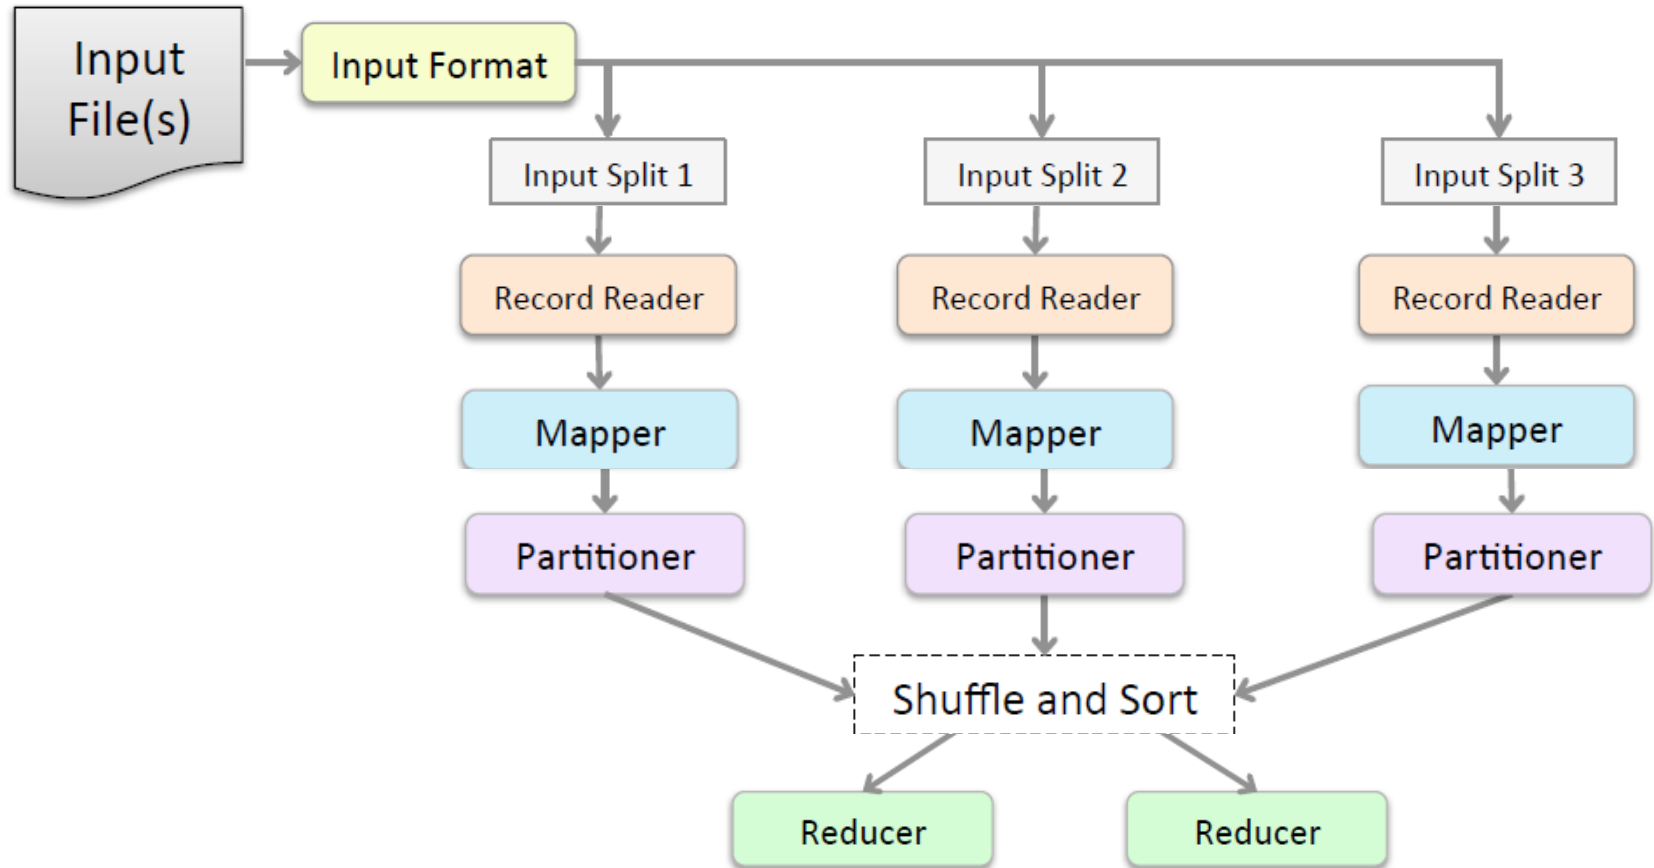
The MapReduce Flow



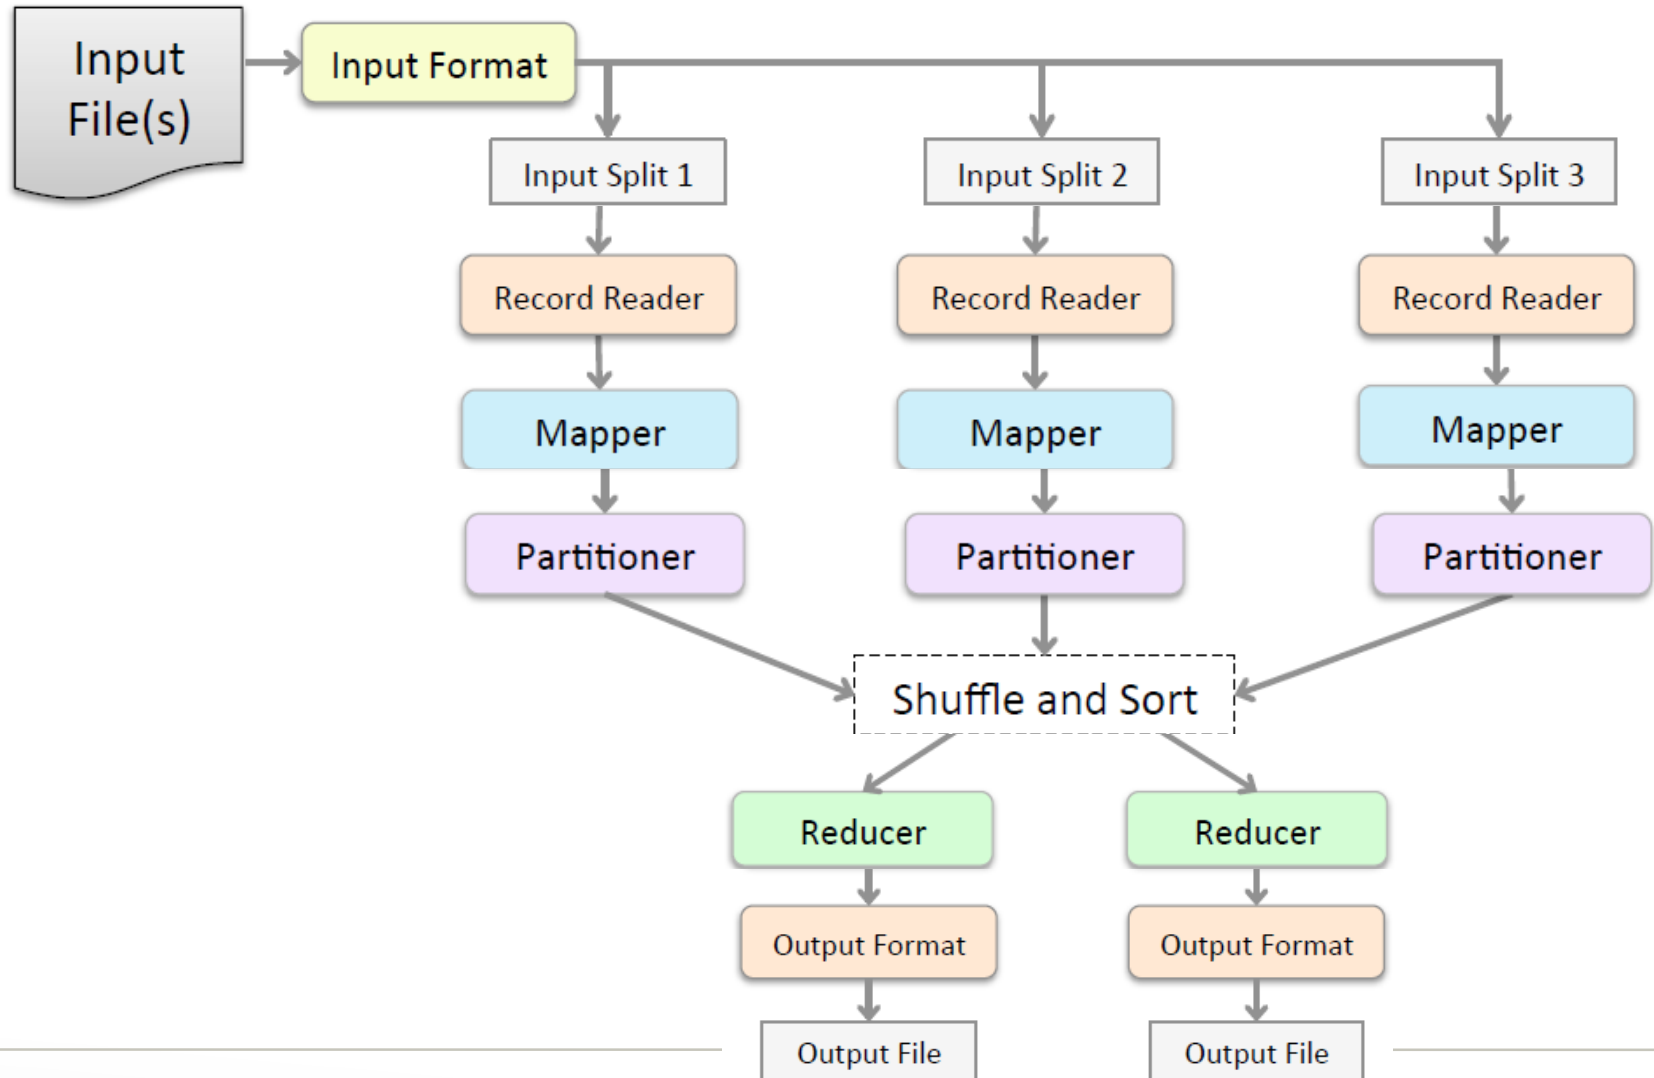
The MapReduce Flow



The MapReduce Flow



The MapReduce Flow



Word Count Example (1)

Input Data

```
the cat sat on the mat  
the aardvark sat on the sofa
```

Map

Reduce

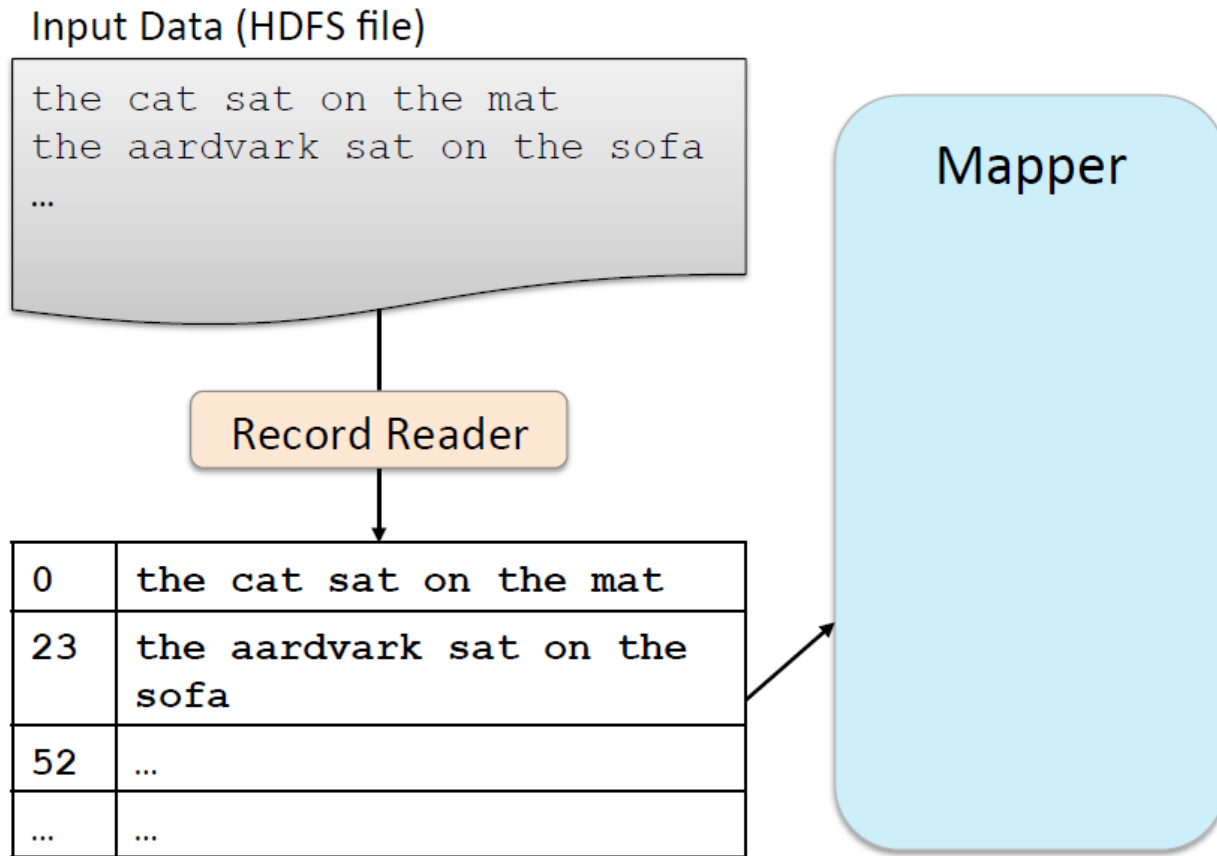
Result

aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4

Word Count Example (2)

- Mapper
 - Input: value: lines of text of input
 - Output: key: word, value: 1
- Reducer
 - Input: key: word, value: set of counts
 - Output: key: word, value: sum
- Launching program
 - Defines this job
 - Submits job to cluster

Word Count Example (3)



Word Count Example (4)

Input Data (HDFS file)

```
the cat sat on the mat
the aardvark sat on the sofa
...
```

Record Reader

0	the cat sat on the mat
23	the aardvark sat on the sofa
52	...
...	...

Mapper

map()

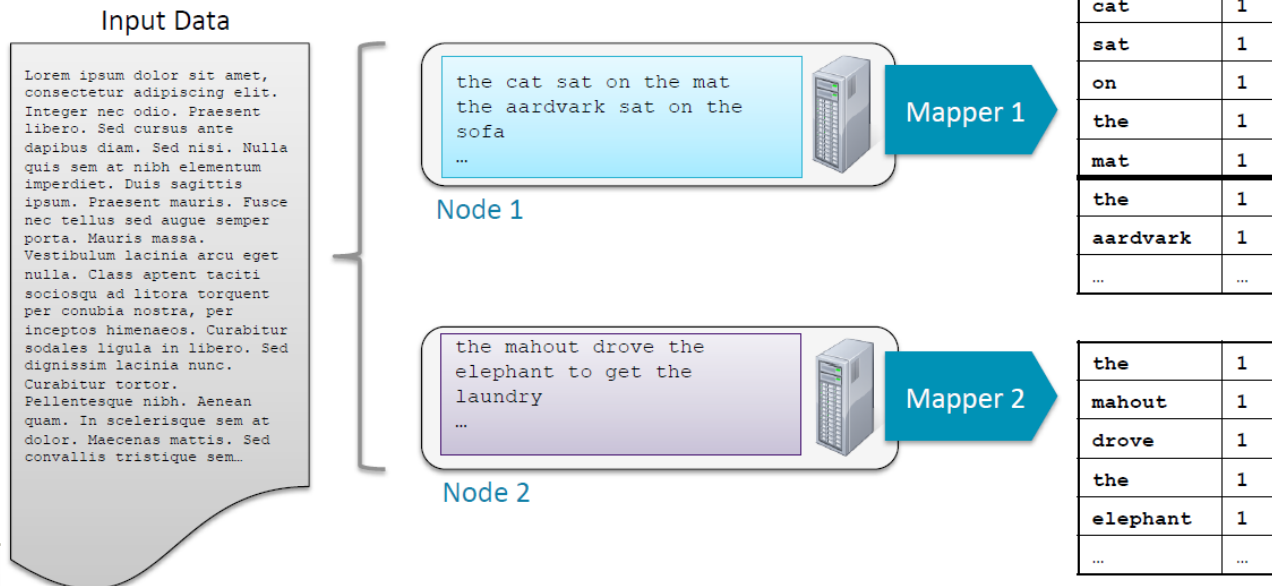
map()

the	1
cat	1
sat	1
on	1
the	1
mat	1

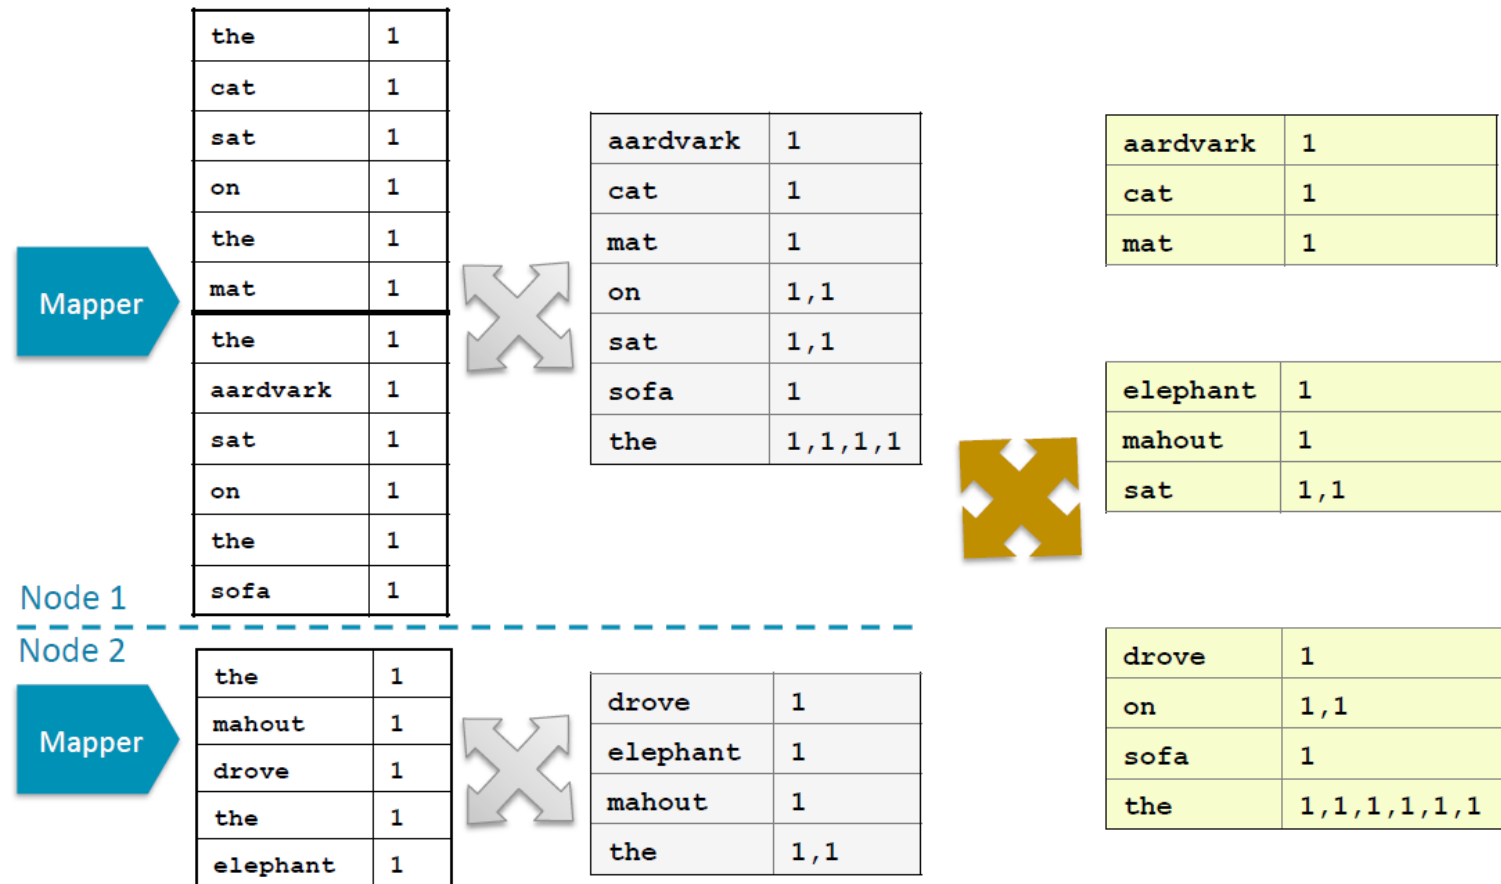
the	1
aardvark	1
sat	1
on	1
the	1
sofa	1

Word Count Example (5)

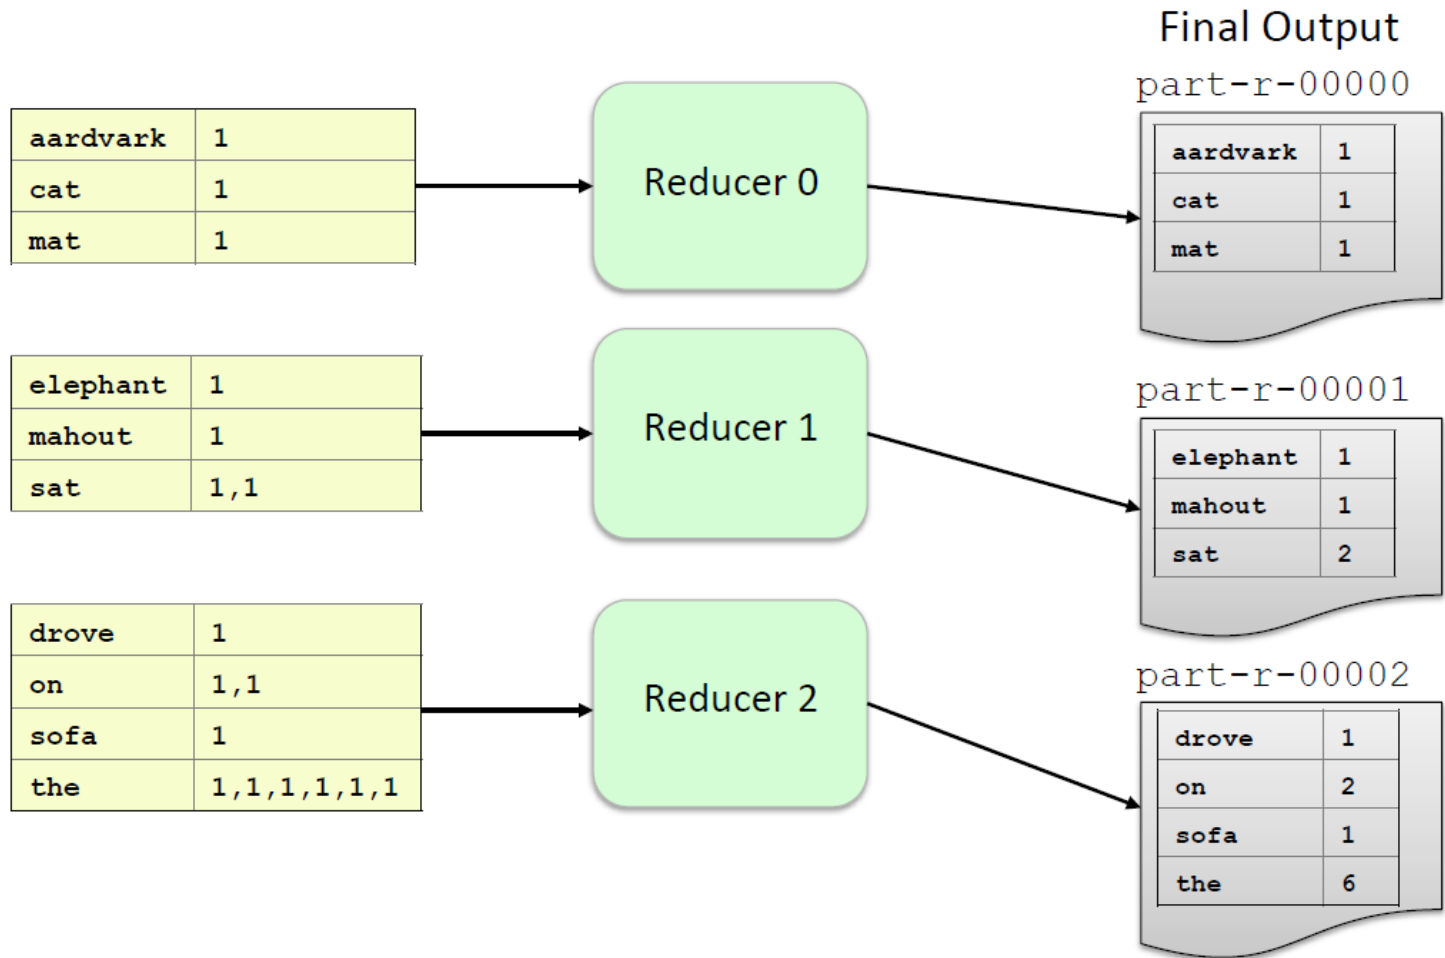
- Hadoop runs Map tasks on the node storing the data (when possible)
 - Minimizes network traffic
 - Many Mappers can run in parallel



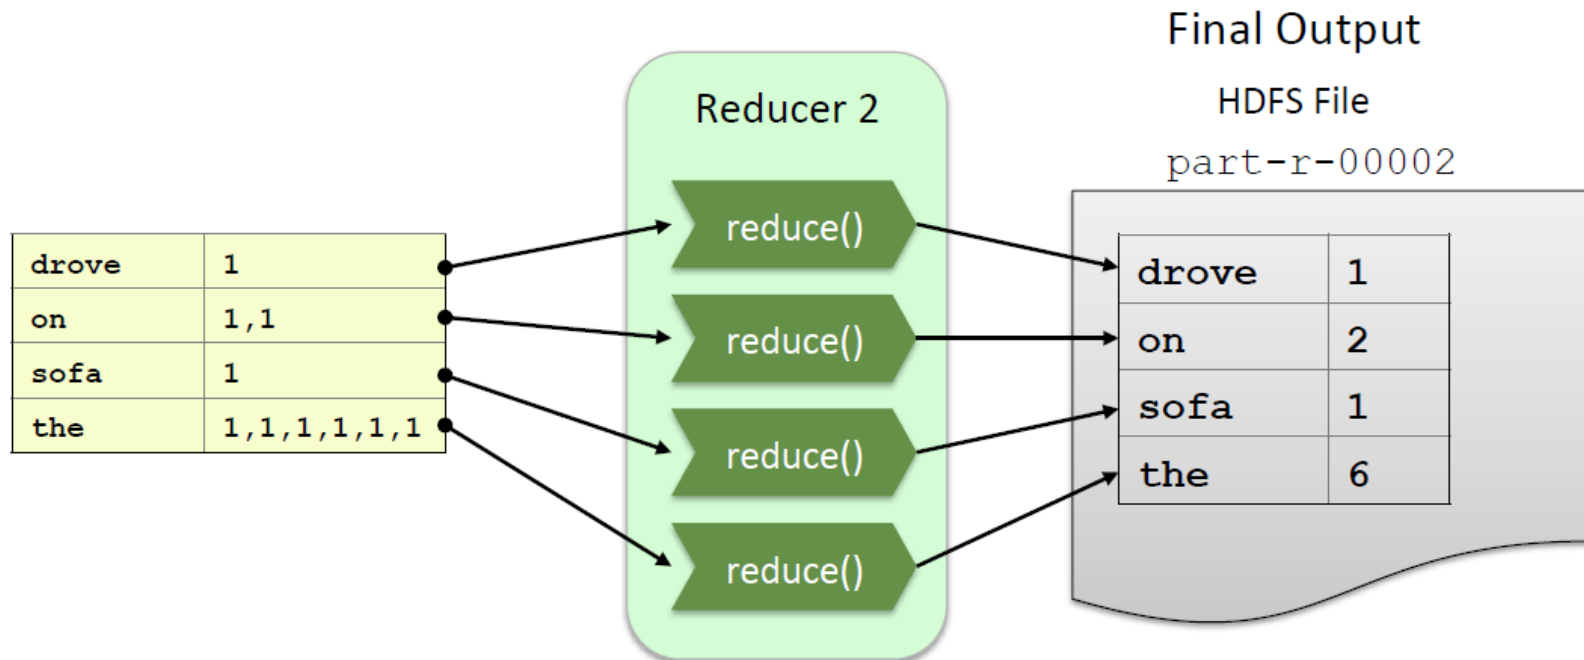
Shuffle and Sort



Example: SumReducer (1)

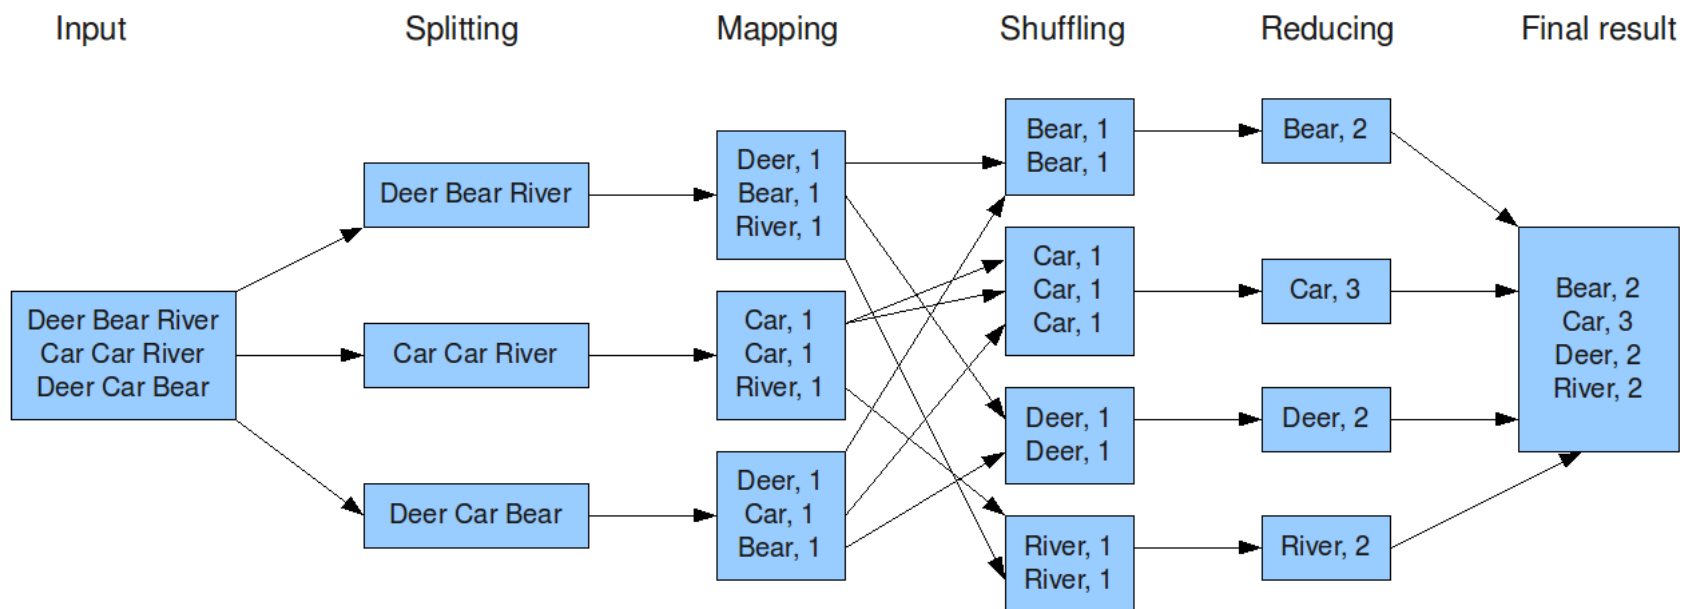


Example: SumReducer (2)



The Final Word Count Dataflow

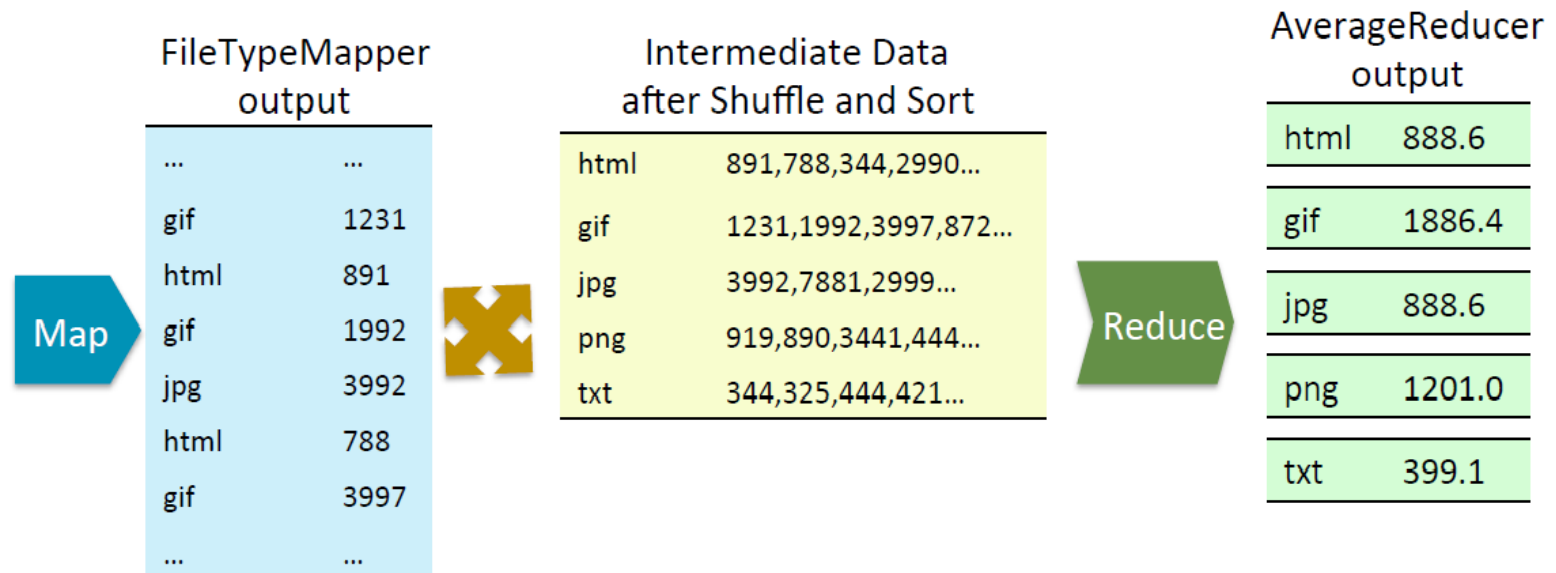
The overall MapReduce word count process



Another Example

Input Data

```
...  
2013-03-15 12:39 - 74.125.226.230 /common/logo.gif 1231ms - 2326  
2013-03-15 12:39 - 157.166.255.18 /catalog/cat1.html 891ms - 1211  
2013-03-15 12:40 - 65.50.196.141 /common/logo.gif 1992ms - 1198  
2013-03-15 12:41 - 64.69.4.150 /common/promoex.jpg 3992ms - 2326  
...
```



Word Count Mapper

```
public static class Map extends MapReduceBase implements
    Mapper<LongWritable, Text, Text, IntWritable> {
    private static final IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public static void map(LongWritable key, Text value,
        OutputCollector<Text, IntWritable> output, Reporter reporter) throws
        IOException {
        String line = value.toString();
        StringTokenizer = new StringTokenizer(line);
        while(tokenizer.hasNext()) {
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

MapReduce: The Mapper (1)

- The Mapper
 - Input: key/value pair
 - Output: A list of zero or more key value pairs

```
map(in_key, in_value) →  
      (inter_key, inter_value) list
```

<i>input key</i>	<i>input value</i>
----------------------	------------------------



<i>intermediate key 1</i>	<i>value 1</i>
<i>intermediate key 2</i>	<i>value 2</i>
<i>intermediate key 3</i>	<i>value 3</i>
...	...

MapReduce: The Mapper (2)

- Input: one line of a file at a time in key/value format
 - The key is the byte offset into the file at which the line starts
 - The value is the contents of the line itself
- Output: in the forms of key/value pairs
 - Input: key/value pair
 - Output: A list of zero or more key value pairs

23	the aardvark sat on the sofa
----	------------------------------

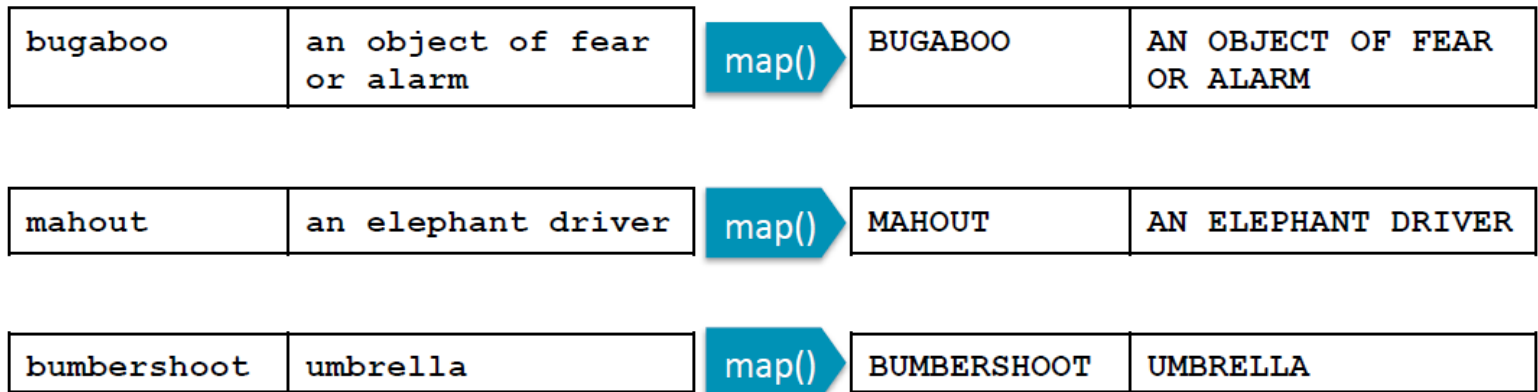


the	1
aardvark	1
sat	1
on	1
the	1
sofa	1

Example Mapper: Upper Case Mapper

- Turn input into upper case (pseudo-code)

```
let map(k, v) =  
    emit(k.toUpper(), v.toUpper())
```



Example Mapper: 'Explode' Mapper

- Output each input character separately (pseudo-code)

```
let map(k, v) =  
  foreach char c in v:  
    emit (k, c)
```

pi	3.14
----	------

map()

pi	3
pi	.
pi	1
pi	4

145	kale
-----	------

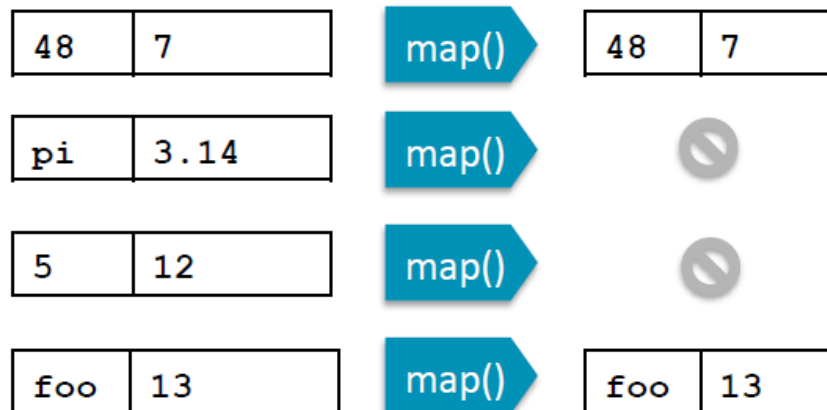
map()

145	k
145	a
145	l
145	e

Example Mapper: 'Filter' Mapper

- Only output key/value pairs where the input value is a prime number (pseudo-code)

```
let map(k, v) =  
  if (isPrime(v)) then emit(k, v)
```



Example Mapper: Changing Keyspaces

- Output the word length as the key (pseudo-code)

```
let map(k, v) =  
    emit(v.length(), v)
```

001	hadoop
-----	--------

map()

6	hadoop
---	--------

002	aim
-----	-----

map()

3	aim
---	-----

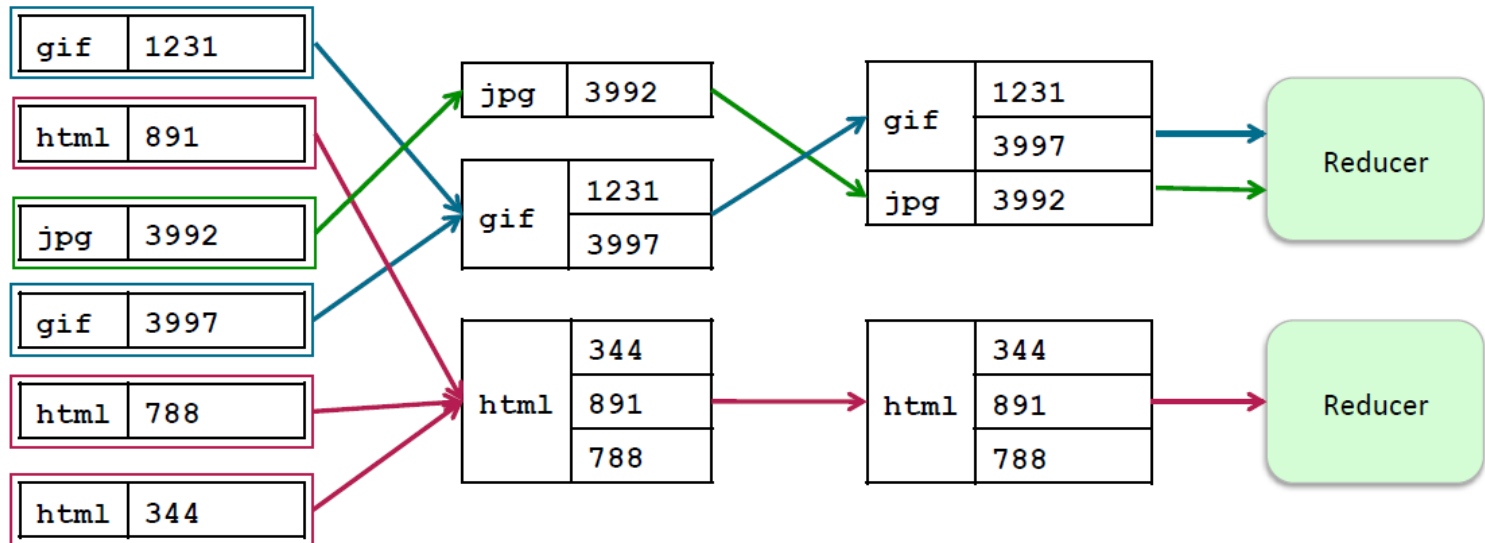
003	ridiculous
-----	------------

map()

10	ridiculous
----	------------

Shuffle and Sort

- After the Map phase is over, all intermediate values for a given intermediate key are grouped together
- Each key and value list is passed to Reducer
 - All values for a particular intermediate key go to the same Reducer
 - The intermediate key/value lists are passed in sorted key order



Word Count Reducer


```
public static class Reduce extends MapReduceBase implements  
    Reducer<Text,IntWritable,Text,IntWritable> {  
public static void map(Text key, Iterator<IntWritable> values,  
    OutputCollector<Text,IntWritable> output, Reporter reporter) throws  
IOException {  
    int sum = 0;  
    while(values.hasNext()) {  
        sum += values.next().get();  
    }  
    output.collect(key, new IntWritable(sum));  
}  
}
```

The Reducer

- The Reducer outputs zero or more final key/value pairs
 - In practice, usually emits a single key/value pair for each input key
 - These are written into HDFS


```
reduce(inter_key, [v1, v2, ...]) →  
    (result_key, result_value)
```

gif	1231
	3997



gif	2614
-----	------

html	344
	891
	788




html	1498
------	------

Example Reducer: Sum Reducer

- All up all the values associated with each intermediate key (pseudo-code)


```
let reduce(k, vals) =  
    sum = 0  
    foreach int i in vals:  
        sum += i  
    emit(k, sum)
```

the	1
	1
	1
	1



the	4
-----	---

SKU0021	34
	8
	19




SKU0021	61
---------	----

Example Reducer: Average Reducer

- Find the mean of all the values associated with each intermediate key (pseudo-code)


```
let reduce(k, vals) =  
  sum = 0; counter = 0;  
  foreach int i in vals:  
    sum += i; counter += 1;  
  emit(k, sum/counter)
```

the	1
	1
	1
	1



the	1
-----	---

SKU0021	34
	8
	19




SKU0021	20.33
---------	-------

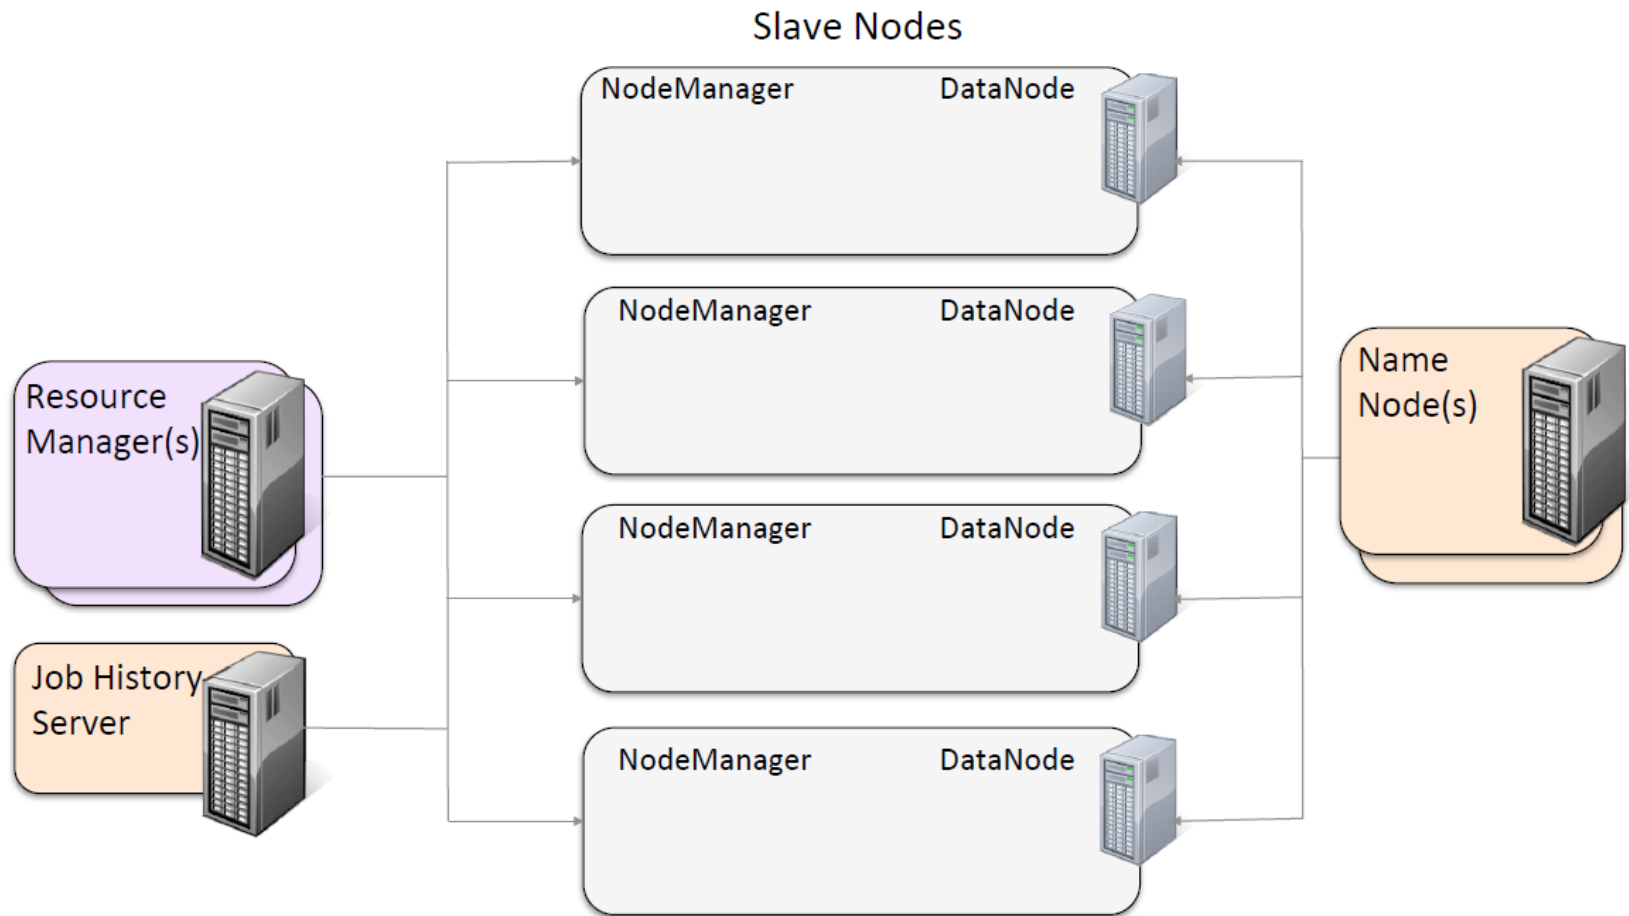
Hadoop Jobs and Tasks

- A **job** is a 'full program'
 - A complete execution of Mappers and Reducers over a dataset
- A **task** is the execution of a single Mapper or Reducer over a slice of data
- A **task attempt** is a particular instance of an **attempt** to execute a task
 - Number of task attempt \geq number of tasks
 - If a task attempt fails, another will be started by the JobTracker or ApplicationMaster
 - **Speculative execution** can also result in more task attempts than completed tasks


Submitting A Job



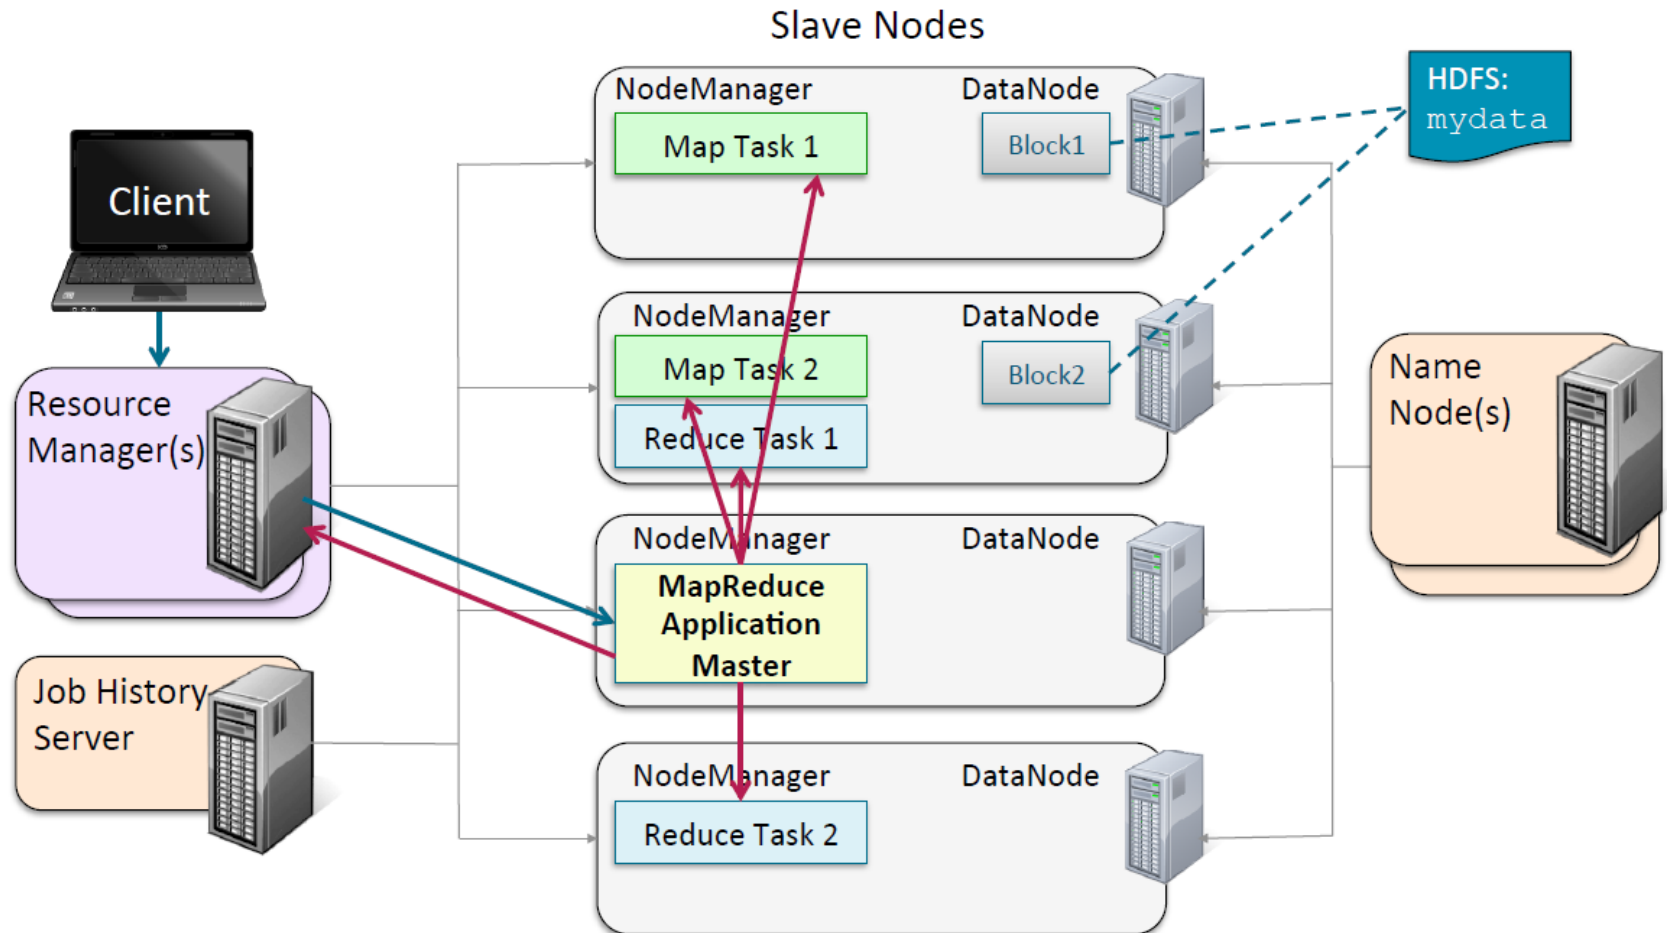
A MapReduce v2 Cluster



Running a Job on a MapReduce v2 Cluster (1)

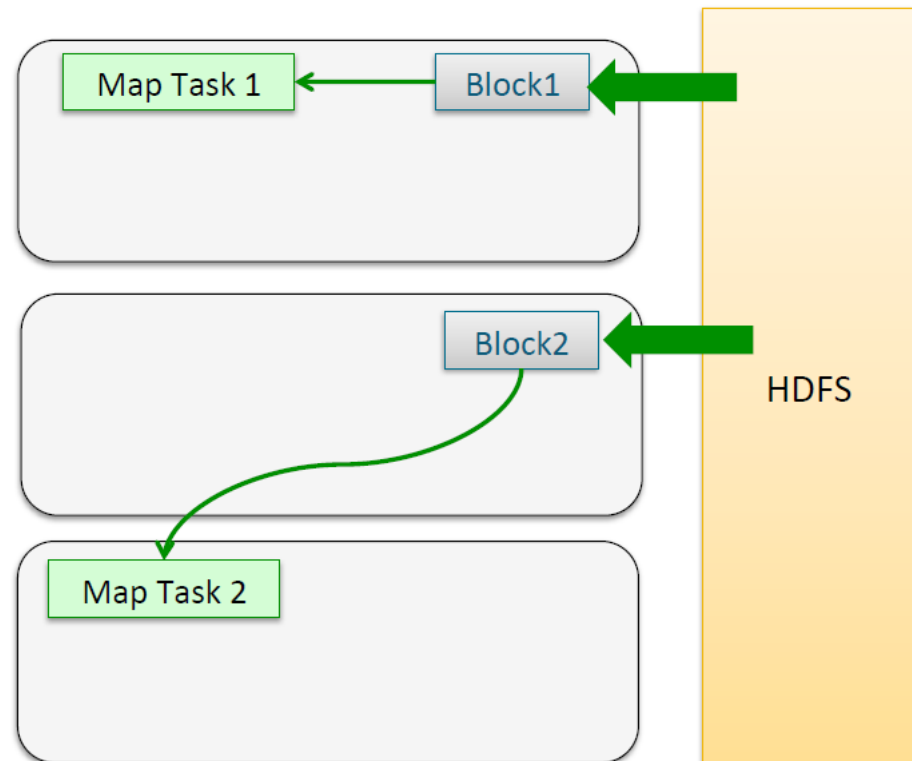


Running a Job on a MapReduce v2 Cluster (2)



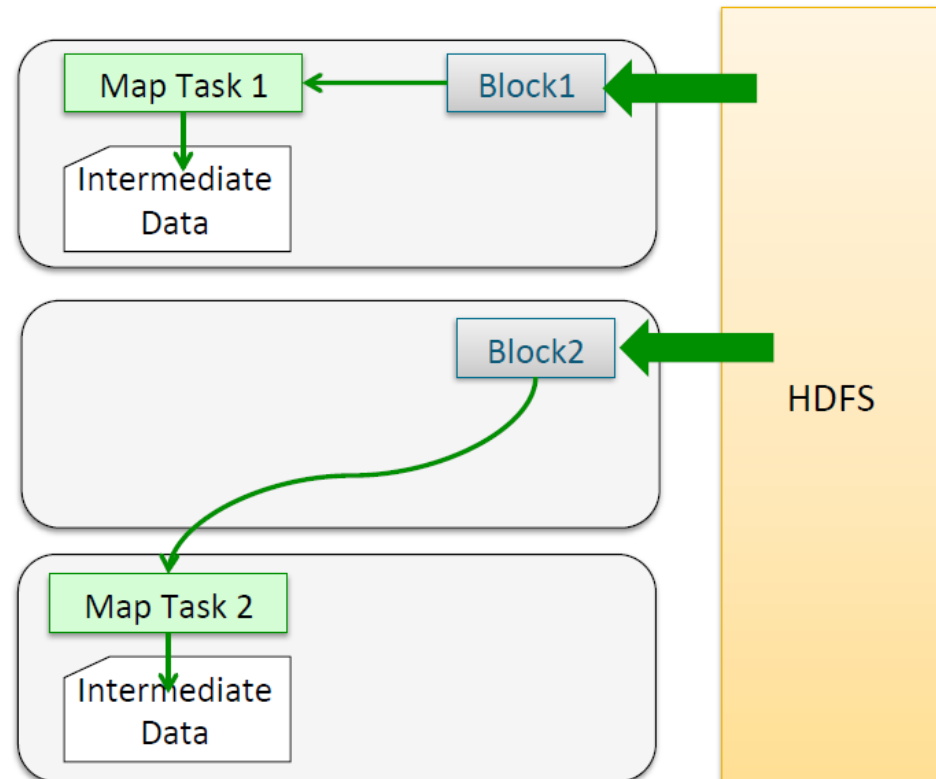
Job Data: Mapper data Locality

- When possible, Map tasks run on a node where a block of data is stored locally
- Otherwise, the Map tasks will transfer the data across the network as it possesses that data



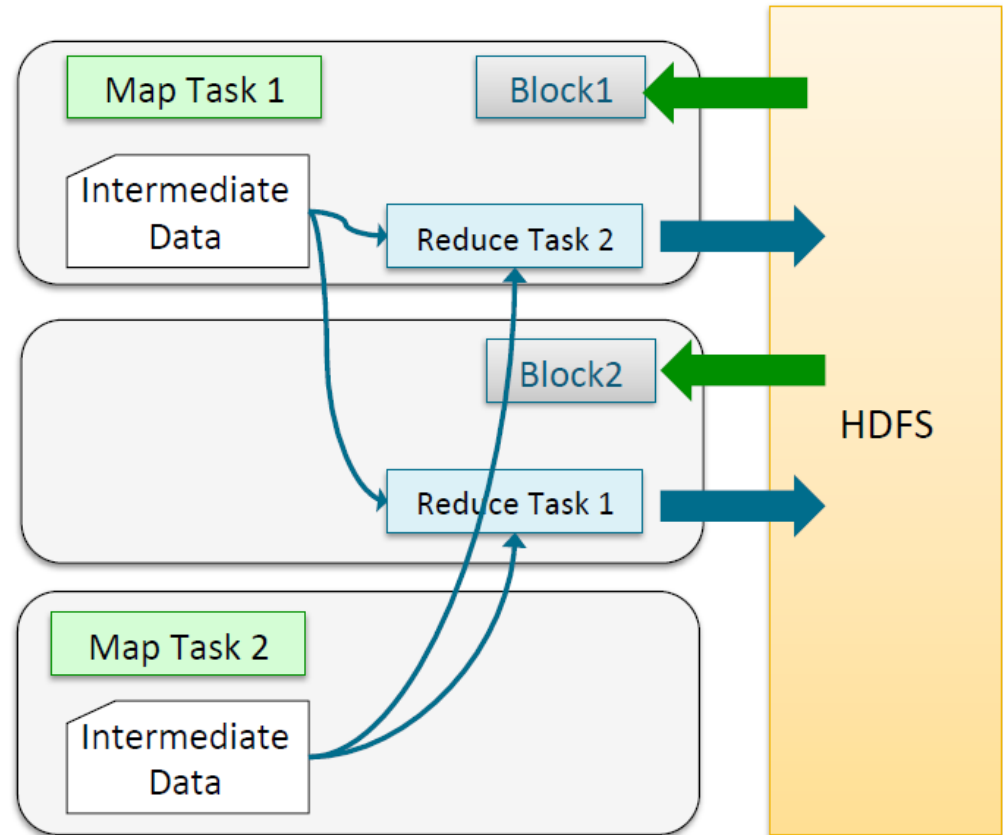
Job Data: Intermediate Data

- The intermediate data of Map tasks is stored on the local disk (not HDFS)



Job Data: Shuffle and Sort

- Intermediate data is transferred across the network to the Reducers
- Reducers write their output to HDFS



Why Shuffle and Sort is a Bottleneck?

- The reduce method in the Reducers cannot start until all Mappers have finished
- In practice, Hadoop will start to transfer and sort data from the Mappers to Reducers as soon as the Mappers finish work
 - This improves the bottleneck a little

Is a Slow Mapper a Bottleneck?

- Faulty hardware or slow machine might degrade the performance
- Hadoop uses **speculative execution** to mitigate against this by creating a new instance of the Mapper to replace the old one, operating on the same data

Creating and Running a MapReduce Job

- Write the Mapper and Reducers classes
- Write a Driver class that configures the job and submits to the cluster (discussed later)
- Compile the Mapper, Reducer, and Driver classes

```
$ javac -classpath `hadoop classpath` MyMapper.java  
MyReducer.java MyDriver.java
```

- Create a jar file with the Mapper, Reducer, and Driver classes

```
$ jar cf MyMR.jar MyMapper.class MyReducer.class  
MyDriver.class
```

- Run the Hadoop jar command to submit the job to the Hadoop cluster

```
$ hadoop jar MyMR.jar MyDriver in_file out_dir
```

To Write a MapReduce Program in Java

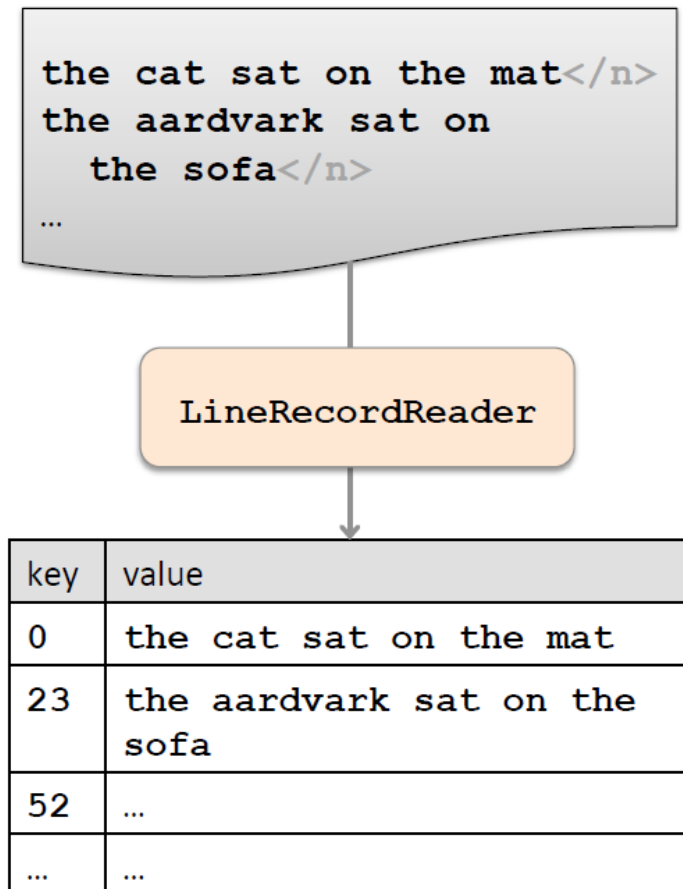
- The driver code
 - Code that runs on the client to configure and submit the job
- The Mapper
- The Reducer

Getting Data to the Mapper

- The data passed to the Mapper is specified by an **InputFormat**
 - Specified in the driver code
 - Defines the location of the input data, a file or a dictionary
 - Determines how to split the input data into **input splits**
 - Each mapper deals with a single input split
 - Creates a **RecordReader** object
 - RecordReader passes the input data into key/value pairs to pass to the Mapper

Example: TextInputFormat

- **TextInputFormat**
 - The default
 - Creates LineRecordReader objects
 - Treats each `\n`-terminated line of a file as a value
 - Key is the byte offset of that line within a file



Other Standard InputFormat

- FileInputFormat
 - Abstract base class for all file-based InputFormats
- KeyValueInputFormat
 - Maps \n-terminated lines as 'key [separator] value'
 - By default, [separator] is a tab
- SequenceFileInputFormat
 - Binary file of (key,value) pairs with some additional metadata
- SequenceFileAsTextInputFormat
 - Similar, but maps (key.toString(), value.toString())

Keys and Values are Objects

- Keys and Values in Hadoop are Java objects
 - Not primitives
- Values are objects which implement **Writable**
- Keys are objects which implement **WritableComparable**

Hadoop Writable Classes

- IntWritable for ints
- LongWritable for longs
- FloatWritable for floats
- DoubleWritable for doubles
- Text for strings
- Etc.

The Driver: Complete Code (1)

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.Job;
public class WordCount {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");
            System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(WordCount.class);
        job.setJobName("Word Count");
    }
}
```

The Driver: Complete Code (2)

```
FileInputFormat.setInputPaths(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));  
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);  
job.setMapOutputKeyClass(Text.class);  
job.setMapOutputValueClass(IntWritable.class);  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);  
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}  
}
```

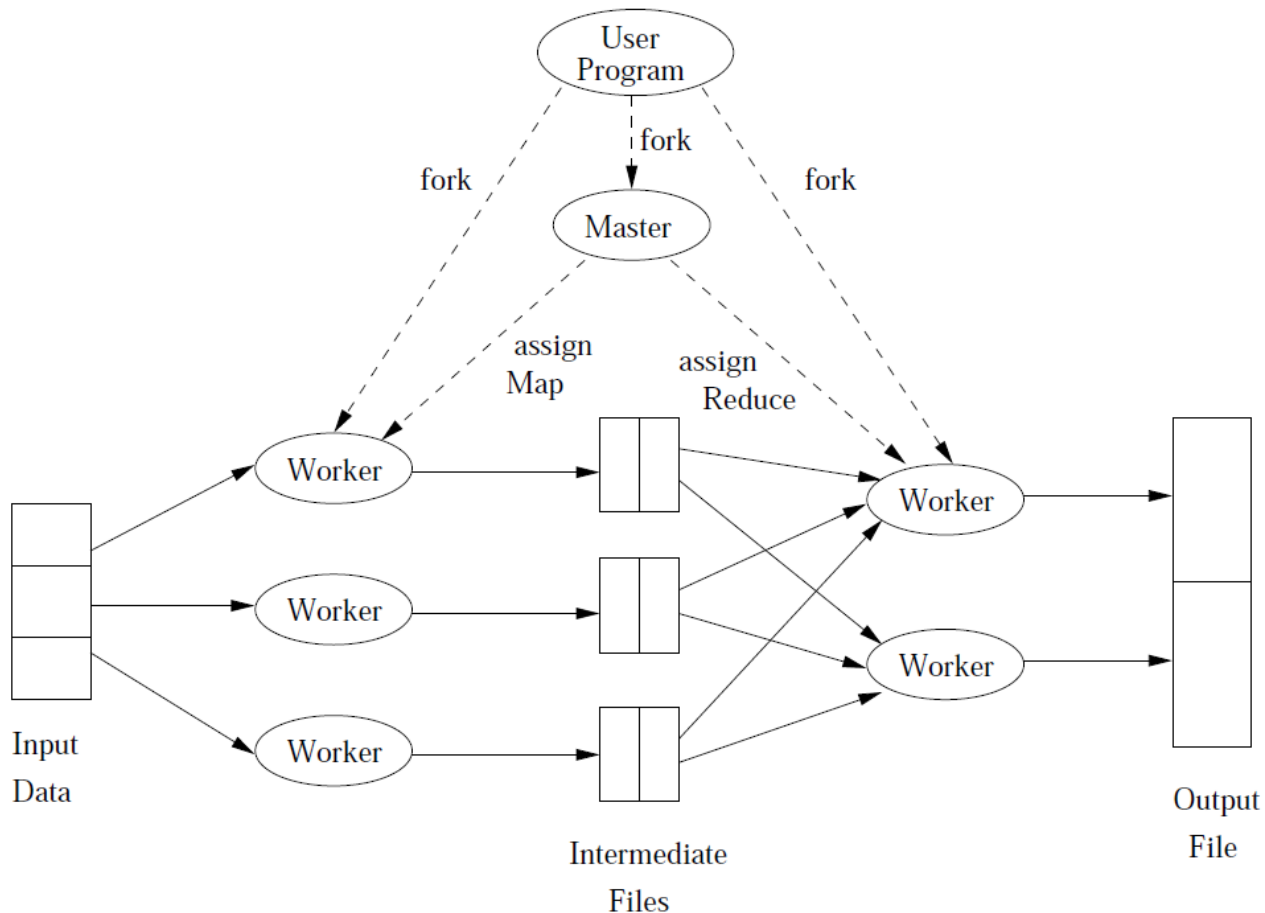
The Mapper Complete Code

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class WordMapper extends Mapper<LongWritable, Text, Text,
IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        for (String word : line.split("\\W+")) {
            if (word.length() > 0) {
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```


The Reducer Complete Code

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class SumReducer extends Reducer<Text, IntWritable, Text,
IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
        context.write(key, new IntWritable(wordCount));
    }
}
```

Execution of a MapReduce Program



Writing Hadoop Code Using Eclipse

- Eclipse is an integrated development environment (IDE)
 - Open source
 - Very popular among Java developers
 - Has plug-ins to speed development in several different languages

Homework3: Writing MapReduce Programs

- Writing a MapReduce Java Program
 - Average Word Length
- More Practice with MapReduce Java programs
 - Log File Analysis
- Writing a report to describe your work

Any Questions?