

Big Data Analysis Platforms



SHYI-CHYI CHENG

Slides credited to Matei Zaharia, UC Berkeley

Outline

- Review of Virtual Machine (虛擬機器回顧)
- Hadoop Platform (運算分析系統架構)
- MapReduce
- Introduction to Python (Python入門簡介)
- Python Spark Platform (Python Spark運算分析架構)
- **Parallel Programming With Spark**

Why Spark?

- Fast, expressive cluster computing system compatible with Apache Hadoop
 - Works with any Hadoop-supported storage system (HDFS, S3, Avro, ...)
- Improves **efficiency** through:
 - In-memory computing primitives
 - General computation graphs  Up to 100× faster
- Improves **usability** through:
 - Rich APIs in Java, Scala, Python
 - Interactive shell  Often 2-10× less code

How to Run It

- Local multicore: just a library in your program
- EC2: scripts for launching a Spark cluster
- Private cluster: Mesos, YARN, Standalone Mode

Languages

- APIs in Java, Scala and Python
- Interactive shells in Scala and Python

Key Idea

- **Work with distributed collections as you would with local ones**
- Concept: resilient distributed datasets (RDDs)
 - Immutable collections of objects spread across a cluster
 - Built through parallel transformations (map, filter, etc)
 - Automatically rebuilt on failure
 - Controllable persistence (e.g. caching in RAM)

Operations

- Transformations (e.g. map, filter, groupBy, join)
 - Lazy operations to build RDDs from other RDDs
- Actions (e.g. count, collect, save)
 - Return a result or write it to storage

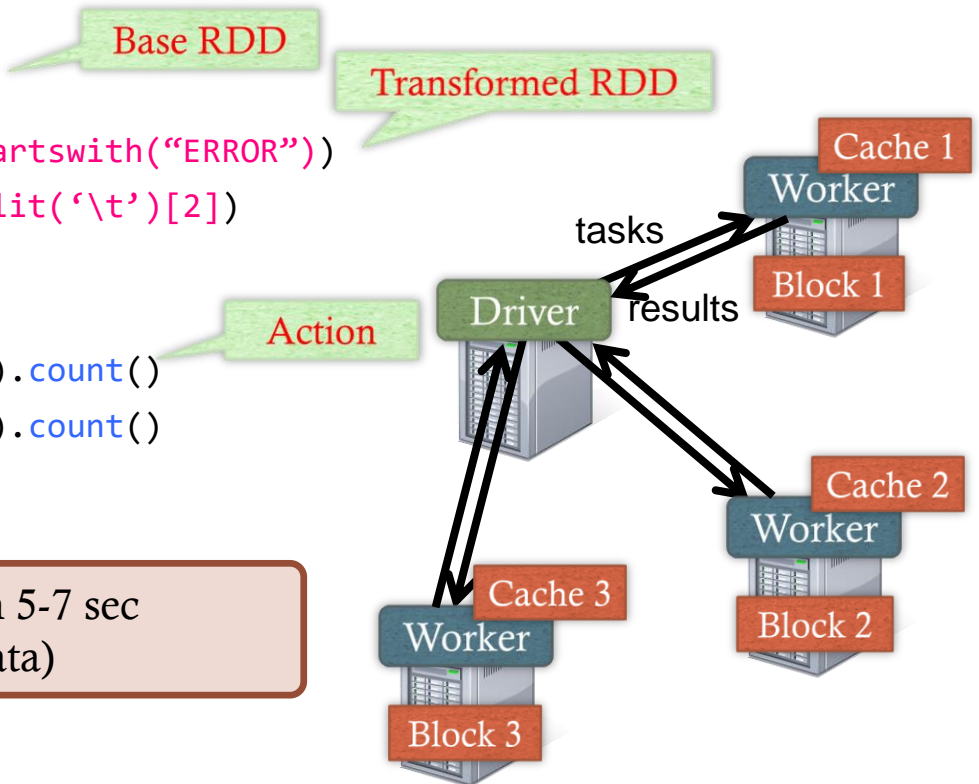
Example: Mining Console Logs

- Load error messages from a log into memory, then interactively search for patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split('\t')[2])
messages.cache()

messages.filter(lambda s: "foo" in s).count()
messages.filter(lambda s: "bar" in s).count()
...
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)



RDD Fault Tolerance


RDDs track the transformations used to build them (their *lineage*) to recompute lost data

E.g:


```
messages = textFile(...).filter(lambda s: s.contains("ERROR"))  
                                .map(lambda s: s.split('\t')[2])
```



Fault Recovery Test



Behavior with Less RAM



Spark in Java and Scala

Java API:

```
JavaRDD<String> lines =  
spark.textFile(...);
```

```
errors = lines.filter(  
    new Function<String,  
    Boolean>() {  
        public Boolean  
        call(String s) {  
            return  
            s.contains("ERROR");  
        }  
    });
```

```
errors.count()
```

Scala API:

```
val lines =  
spark.textFile(...)
```

```
errors = lines.filter(s =>  
    s.contains("ERROR"))  
// can also write  
filter(_.contains("ERROR"))
```

```
errors.count
```

Which Language Should I Use?

- Standalone programs can be written in any, but console is only Python & Scala
- **Python developers:** can stay with Python for both
- **Java developers:** consider using Scala for console (to learn the API)
- Performance: Java / Scala will be faster (statically typed), but Python can do well for numerical work with NumPy

Scala Cheat Sheet

More details:
scala-lang.org

Variables:

```
var x: Int = 7
var x = 7    // type
              inferred
val y = "hi"  // read-
              only
```

Functions:

```
def square(x: Int): Int
= x*x

def square(x: Int): Int
= {
  x*x    // last line
  returned
}
```

Collections and closures:

```
val nums = Array(1, 2, 3)
nums.map((x: Int) => x + 2) // =>
Array(3, 4, 5)
nums.map(x => x + 2)       // => same
nums.map(_ + 2)           // => same
nums.reduce((x, y) => x + y) // => 6
nums.reduce(_ + _)         // => 6
```

Java interop:

```
import java.net.URL
new
URL("http://cnn.com").openStream()
```

Learning Spark

- Easiest way: Spark interpreter (spark-shell or pyspark)
 - Special Scala and Python consoles for cluster use
- Runs in local mode on 1 thread by default, but can control with MASTER environment var:

```
MASTER=local      ./spark-shell      # local, 1 thread
MASTER=local[2]   ./spark-shell      # local, 2
threads
MASTER=spark://host:port ./spark-shell # Spark
standalone cluster
```

First Stop: SparkContext

- Main entry point to Spark functionality
- Created for you in Spark shells as variable `sc`
- In standalone programs, you'd make your own (see later for details)

Creating RDDs

Turn a local collection into an RDD

```
sc.parallelize([1, 2, 3])
```

Load text file from local FS, HDFS, or S3

```
sc.textFile("file.txt")
```

```
sc.textFile("directory/*.txt")
```

```
sc.textFile("hdfs://namenode:9000/path/file")
```

Use any existing Hadoop InputFormat

```
sc.hadoopFile(keyClass, valClass, inputFmt, conf)
```

Basic Transformations

```
nums = sc.parallelize([1, 2, 3])

# Pass each element through a function
squares = nums.map(lambda x: x*x)    # => {1, 4, 9}

# Keep elements passing a predicate
even = squares.filter(lambda x: x % 2 == 0) # => {4}

# Map each element to zero or more others
nums.flatMap(lambda x: range(0, x))  # => {0, 0, 1, 0, 1, 2}
```

Range object (sequence
of numbers 0, 1, ..., x-1)

Basic Actions

```
nums = sc.parallelize([1, 2, 3])

# Retrieve RDD contents as a local collection
nums.collect() # => [1, 2, 3]

# Return first K elements
nums.take(2)    # => [1, 2]

# Count number of elements
nums.count()    # => 3

# Merge elements with an associative function
nums.reduce(lambda x, y: x + y) # => 6

# Write elements to a text file
nums.saveAsTextFile("hdfs://file.txt")
```

Working with Key-Value Pairs

- Spark's “distributed reduce” transformations act on RDDs of *key-value pairs*
- Python:

```
pair = (a, b)
```

```
pair[0] # => a
```

```
pair[1] # => b
```
- Scala:

```
val pair = (a, b)
```

```
pair._1 // => a
```

```
pair._2 // => b
```
- Java:

```
Tuple2 pair = new Tuple2(a, b); // class
```

```
scala.Tuple2
```

```
pair._1 // => a
```

```
pair._2 // => b
```

Some Key-Value Operations

```
pets = sc.parallelize([("cat", 1), ("dog", 1), ("cat", 2)])
```

```
pets.reduceByKey(lambda x, y: x + y)
```

```
# => {(cat, 3), (dog, 1)}
```

```
pets.groupByKey()
```

```
# => {(cat, Seq(1, 2)), (dog, Seq(1))}
```

```
pets.sortByKey()
```


```
# => {(cat, 1), (cat, 2), (dog, 1)}
```

reduceByKey also automatically implements combiners on the map side

Example: Word Count

```
lines = sc.textFile("hamlet.txt")
```

```
counts = lines.flatMap(lambda line: line.split(" ")) \
                .map(lambda word: (word, 1)) \
                .reduceByKey(lambda x, y: x + y)
```



Multiple Datasets

```
visits = sc.parallelize([(“index.html”, “1.2.3.4”),  
                        ( “about.html”, “3.4.5.6”),  
                        ( “index.html”, “1.3.3.1”)])
```

```
pageNames = sc.parallelize([(“index.html”, “Home”),  
                           ( “about.html”, “About”)])
```

```
visits.join(pageNames)
```

```
# ( “index.html”, ( “1.2.3.4”, “Home” ))
```

```
# ( “index.html”, ( “1.3.3.1”, “Home” ))
```

```
# ( “about.html”, ( “3.4.5.6”, “About” ))
```

```
visits.cogroup(pageNames)
```

```
# ( “index.html”, (Seq(“1.2.3.4”, “1.3.3.1”), Seq(“Home”)))
```

```
# ( “about.html”, (Seq(“3.4.5.6”), Seq(“About”)))
```

Controlling the Level of Parallelism

- All the pair RDD operations take an optional second parameter for number of tasks

```
words.reduceByKey(lambda x, y: x + y,  
5)
```

```
words.groupByKey(5)
```

```
visits.join(pageViews, 5)
```


Using Local Variables

- External variables you use in a closure will automatically be shipped to the cluster:

```
query = raw_input("Enter a query:")  
pages.filter(lambda x:  
x.startswith(query)).count()
```

- Some caveats:
 - Each task gets a new copy (updates aren't sent back)
 - Variable must be Serializable (Java/Scala) or Pickle-able (Python)
 - Don't use fields of an outer object (ships all of it!)

Closure Mishap Example

```
class MyCoolRddApp {  
  val param = 3.14  
  val log = new Log(...)  
  ...  
  
  def work(rdd: RDD[Int]) {  
    rdd.map(x => x + param)  
        .reduce(...)   
  }  
}
```

NotSerializableException:
MyCoolRddApp (or Log)

How to get around it:

```
class MyCoolRddApp {  
  ...  
  
  def work(rdd: RDD[Int]) {  
    val param_ = param  
    rdd.map(x => x + param_)  
        .reduce(...)   
  }  
}
```


References only local
variable instead of
this.param

More Details

- Spark supports lots of other operations!
- Full programming guide: spark-project.org/documentation


Software Components

- Spark runs as a library in your program
(one instance per app)
- Runs tasks locally or on a cluster
 - Standalone deploy cluster, Mesos or YARN
- Accesses storage via Hadoop InputFormat API
 - Can use HBase, HDFS, S3, ...



Task Scheduler

- Supports general task graphs
- Pipelines functions where possible
- Cache-aware data reuse & locality
- Partitioning-aware to avoid shuffles



Hadoop Compatibility

- Spark can read/write to any storage system / format that has a plugin for Hadoop!
 - Examples: HDFS, S3, HBase, Cassandra, Avro, SequenceFile
 - Reuses Hadoop's InputFormat and OutputFormat APIs
- APIs like `SparkContext.textFile` support filesystems, while `SparkContext.hadoopRDD` allows passing any Hadoop JobConf to configure an input source

Build Spark

- Requires Java 6+, Scala 2.9.2

```
git clone git://github.com/mesos/spark
cd spark
sbt/sbt package
```

```
# Optional: publish to local Maven cache
sbt/sbt publish-local
```

Add Spark to Your Project

- Scala and Java: add a Maven dependency on

groupId: `org.spark-project`

artifactId: `spark-core_2.9.1`

version: `0.7.0-SNAPSHOT`

- Python: run program with our pyspark script

Create a SparkContext

Scala

```
import spark.SparkContext
import spark.SparkContext._

val sc = new SparkContext("masterUrl", "name", "sparkHome", Seq("app.jar"))
```

Cluster URL, or
local / local[N]

App
name

Spark install
path on
cluster

List of JARs
with app code
(to ship)

Java

```
import spark.api.java.JavaSparkContext;

JavaSparkContext sc = new JavaSparkContext(
    "masterUrl", "name", "sparkHome", new String[] {"app.jar"});
```

Python

```
from pyspark import SparkContext

sc = SparkContext("masterUrl", "name", "sparkHome", ["library.py"])
```

Complete App: Scala

```
import spark.SparkContext
import spark.SparkContext._

object WordCount {
  def main(args: Array[String]) {
    val sc = new SparkContext("local", "WordCount", args(0),
Seq(args(1)))
    val lines = sc.textFile(args(2))
    lines.flatMap(_.split(" "))
      .map(word => (word, 1))
      .reduceByKey(_ + _)
      .saveAsTextFile(args(3))
  }
}
```

Complete App: Python

```
import sys
from pyspark import SparkContext

if __name__ == "__main__":
    sc = SparkContext( "local", "WordCount", sys.argv[0], None)
    lines = sc.textFile(sys.argv[1])

    lines.flatMap(lambda s: s.split(" ")) \
        .map(lambda word: (word, 1)) \
        .reduceByKey(lambda x, y: x + y) \
        .saveAsTextFile(sys.argv[2])
```

Local Mode

- Just pass `local` or `local[k]` as master URL
- Still serializes tasks to catch marshaling errors
- Debug using local debuggers
 - For Java and Scala, just run your main program in a debugger
 - For Python, use an attachable debugger (e.g. PyDev, winpdb)
- Great for unit testing

Private Cluster

- Can run with one of:
 - Standalone deploy mode (similar to Hadoop cluster scripts)
 - Apache Mesos: spark-project.org/docs/latest/running-on-mesos.html
 - Hadoop YARN: spark-project.org/docs/0.6.0/running-on-yarn.html
- Basically requires configuring a list of workers, running launch scripts, and passing a special cluster URL to SparkContext

Amazon EC2

- Easiest way to launch a Spark cluster

```
git clone git://github.com/mesos/spark.git
```

```
cd spark/ec2
```

```
./spark-ec2 -k keypair -i id_rsa.pem -s slaves \  
[launch|stop|start|destroy] clusterName
```

- Details: spark-project.org/docs/latest/ec2-scripts.html
- **New: run Spark on Elastic MapReduce – tinyurl.com/spark-emr**

Viewing Logs

- Click through the web UI at `master:8080`
- Or, look at `stdout` and `stderr` files in the Spark or Mesos “work” directory for your app:
 `work/<ApplicationID>/<ExecutorID>/stdout`
- Application ID (Framework ID in Mesos) is printed when Spark connects

Homework5: Work Count in Python Spark

- 下載安裝eclipse Scala IDE
- 安裝pyDev, 並設定環境變數
- 新增pyDev專案, 使用Python Spark重新製作 Word Count的例子
- 紀錄執行結果
- 報告撰寫及繳交

Any Questions?