

Big Data Analytics

Link Analysis

SHYI-CHYI CHENG

Outline

- **Introduction**
- PageRank
- Topic-Sensitive PageRank
- Link Spam
- HITS
- Summary

Motivation

ted演講 - Yahoo 奇摩

https://tw.search.yahoo.com/search?p=ted演講&fr=yfp-s&ei=utf-8&v=0

是新用戶嗎？請註冊 | 登入

信箱 | Yahoo 奇摩

YAHOO! 奇摩

網頁
知識+
圖片
影片
部落格
購物
在地生活
字典
更多

台灣網頁優先
全球網頁
繁體中文

不限時間
過去一天
過去一週

ted演講 搜尋

選項

相關詞： [ted演講中文](#), [ted演講 youtube](#), [ted演講中文字幕](#), [更多...](#)

ted演講 相關廣告

職場必修英文免費演講限額參加 | tutorabc.com

TutorABC線上真人英文講堂，一次報名享十堂課程，限額百名免費參加！
www.tutorabc.com

Browse Talks | TED.com - TED: Ideas worth spreading

TED is a nonprofit devoted to Ideas Worth Spreading - through TED.com, our annual conferences, the annual TED Prize and local TEDx events. ... Please enter an email address. Please enter a valid email address. Did you mean? Please check Daily or Weekly and ...
www.ted.com/talks/browse?language=zh-tw 庫存頁面 - 更多此站結果

TEDxTaipei official website | TEDxTaipei

演講影片 關於TED/TEDx 關於TEDxTaipei TEDxTaipei 團隊 人才招聘 贊助夥伴 聯絡我們
最新文章 What Matters Now大哉問- 10/5前進報導 What Matters Now大哉問- 10/5本質報導
What Matters Now大哉問- 10/5共鳴報導 ...
tedxtaipei.com 庫存頁面 - 更多此站結果

TED Talks 演講 - VoiceTube 《看影片學英語》15,000 部英文學習影片

演講就找 沈芯菱

演講推薦！最震撼人心的演講，百大企業推薦，感動30萬人，絕不可錯過！
www.ananedu.com

驅勢航空講座 免費預約報名

夢想起飛，等你登機！前長榮事務長 Kora經驗分享，台北、桃園場次限額報名！
edm.tlc.tw

更多贊助網站

熱門搜尋


1 最正檢察總長	6 金秀賢胖了
2 6秒賺4個月薪水	7 台大生搶飯吃
3 馬桶要能泡麵	8 頂新商品
4 金正恩神隱原因	9 羅賓威廉斯遺作
5 鍾漢良西遊記	10 金城武極度

How many articles did relate to the user query?

Why we need link analysis?

- The web is **not** just a collection of documents – its **hyperlinks are important**
- **A link from page A to page B** may indicate:
 - A is related to B , or
 - A is recommending, citing, voting for or endorsing B
- Link types
 - **referential** – *click here and get back home*, or
 - **Informational** – *click here to get more detail*
- Links effect the **Page Ranking**

Example Web Graph



Early Search Engines and Term Spam

- Early search engines mainly compare content similarity of the query and the indexed pages
 - IR techniques are applied: cosine, TF-IDF, ...
- **Term spam**
 - Techniques for fooling search engines into believing your page is about something it is not
 - Content similarity is easily spammed

Google Solution

- PageRank was used to simulate where Web surfers
 - Pages that would have **a large number of surfers** were considered more “**important**” than pages that would rarely be visited.
- The content of a page was judged **not only by the terms** appearing on that page, **but by the terms used in or near the links** to that page.

Users of the Web “**vote with their feet.**”

Link Analysis Strategies

- During 1997-1998, two most influential hyperlink based search algorithms **PageRank** and **HITS** were reported.
- Both algorithms are related to **social networks**. They exploit the hyperlinks of the Web to rank pages according to their levels of “prestige” or “authority”.
 - **HITS**: Jon Kleinberg (Cornel University), at *Ninth Annual ACM-SIAM Symposium on Discrete Algorithms*, January 1998
 - **PageRank**: Sergey Brin and Larry Page, PhD students from Stanford University, at *Seventh International World Wide Web Conference (WWW7)* in April, 1998.

PageRank powers the Google search engine.


PageRank

- PageRank is a function that assigns a real number to each page in the Web
- PageRank relies on the democratic nature of the Web by using its vast link structure as an indicator of an individual page's value or quality.
- PageRank interprets a hyperlink from page x to page y as a vote, by page x , for page y .
- However, PageRank looks at more than the sheer number of votes; it also analyzes the page that casts the vote.
 - Votes casted by “important” pages weigh more heavily and help to make other pages more “important.”

More specifically

- A **hyperlink** from a page to another page is an **implicit conveyance of authority** to the target page.
 - The more in-links that a page i receives, the more prestige the page i has.
- **Pages** that point to page i also **have** their own **prestige scores**.
 - A page of a higher prestige pointing to i is more important than a page of a lower prestige pointing to i .
 - In other words, a page is important if it is pointed to by other important pages.

Web as a Directed Graph



$$\begin{array}{c} \begin{array}{c} A \\ B \\ C \\ D \end{array} \begin{bmatrix} A & B & C & D \\ 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \end{array}$$

Transition Matrix M

Probabilities of all Nodes v

- Suppose we start a random surfer at any of the n pages of the Web with equal probability
- Initially, Probability vector $v = [1/n, \dots, 1/n]$
- Loop

$$v = Mv$$

until end the surfing

Probability Estimation by Markov Process

- Conditions
 - The graph is strongly connected; that is, it is possible to get from any node to any other node.
 - There are no dead ends: nodes that have no arcs out.

What is a Markov Chain?

- A Markov chain has two components:
 - 1) A network structure much like a web site, where each node is called a state.
 - 2) A transition probability of traversing a link given that the chain is in a state.
 - For each state the sum of outgoing probabilities is one.
- A sequence of steps through the chain is called a *random walk*.


Markov chains

- Markov chains have been extensively studied by statisticians and have been applied in a wide variety of areas.

Markov Property:

$$P(S_t | S_{t-1}, S_{t-2}, \dots, S_2, S_1) = P(S_t | S_{t-1}) \quad (3.3)$$

Markov Chain Example




The Random Surfer

- Assume the web is a Markov chain.
- Surfers randomly click on links, where the probability of an outlink from page A is $1/m$, where m is the number of outlinks from A .
- The surfer occasionally gets *bored* and is *teleported* to another web page, say B , where B is equally likely to be any page.
- Using the theory of Markov chains it can be shown that if the surfer follows links for long enough, *the PageRank of a web page is the probability that the surfer will visit that page.*

Principal Eigenvector of M

- The **limit** is reached when multiplying the distribution by M another time does not change the distribution.
- The limiting v is an eigenvector of M
- We can compute the principal eigenvector of M by starting with the initial vector v_0 and multiplying by M some number of times, until the vector we get shows little change at each round.

Example



$$M = \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \quad \begin{bmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix} \quad \begin{bmatrix} 15/48 \\ 11/48 \\ 11/48 \\ 11/48 \end{bmatrix} \quad \begin{bmatrix} 11/32 \\ 7/32 \\ 7/32 \\ 7/32 \end{bmatrix} \quad \dots \quad \begin{bmatrix} 3/9 \\ 2/9 \\ 2/9 \\ 2/9 \end{bmatrix}$$

\mathbf{v}

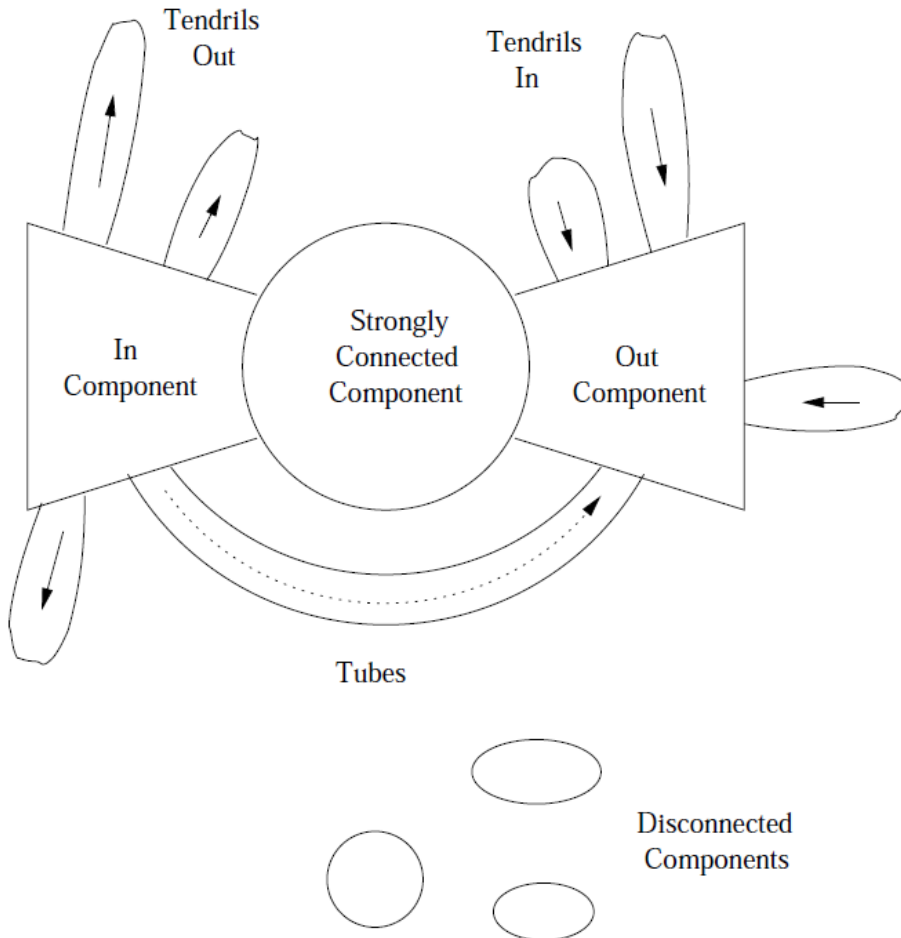
$M\mathbf{v}$

$M^2\mathbf{v}$

$M^3\mathbf{v}$

$M^{56}\mathbf{v}$


Structure of the Web



To use **probability** as measuring the **importance** of a page falsely concludes that nothing in the SCC or in-component is of any importance.

Avoiding Dead Ends

- A page with no link out is called a dead end




$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$


$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 3/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix} \begin{bmatrix} 5/48 \\ 7/48 \\ 7/48 \\ 7/48 \end{bmatrix} \begin{bmatrix} 21/288 \\ 31/288 \\ 31/288 \\ 31/288 \end{bmatrix} \dots \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Two approaches to dealing with dead ends

- We can drop the dead ends from the graph, and also drop their incoming arcs, recursively.
- We can modify the process by which random surfers are assumed to move about the Web.
 - The taxation method



Spider Trap



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 3/24 \\ 5/24 \\ 11/24 \\ 5/24 \end{bmatrix} \begin{bmatrix} 5/48 \\ 7/48 \\ 29/48 \\ 7/48 \end{bmatrix} \begin{bmatrix} 21/288 \\ 31/288 \\ 205/288 \\ 31/288 \end{bmatrix} \cdots \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

Revised Probability Estimation


- The iterative step, where we compute a new vector estimate of PageRanks \mathbf{v}' from the current PageRank estimate \mathbf{v} and the transition matrix M is

$$\mathbf{v}' = \beta M \mathbf{v} + (1 - \beta) \mathbf{e} / n$$

where β is a chosen constant, usually in the range 0.8 to 0.9, \mathbf{e} is a vector of all 1's with the appropriate number of components, and n is the number of nodes in the Web graph.

\mathbf{e} *biases* the jump prefer some pages over others:
--e.g. \mathbf{e} has 1 for your home page and 0 otherwise.
--e.g. \mathbf{e} prefers the topics you are interested in.

Example



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

$$\beta = 0.8$$

$$\mathbf{v}' = \begin{bmatrix} 0 & 2/5 & 0 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 4/5 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{bmatrix} \mathbf{v} + \begin{bmatrix} 1/20 \\ 1/20 \\ 1/20 \\ 1/20 \end{bmatrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \begin{bmatrix} 9/60 \\ 13/60 \\ 25/60 \\ 13/60 \end{bmatrix} \begin{bmatrix} 41/300 \\ 53/300 \\ 153/300 \\ 53/300 \end{bmatrix} \begin{bmatrix} 543/4500 \\ 707/4500 \\ 2543/4500 \\ 707/4500 \end{bmatrix} \dots \begin{bmatrix} 15/148 \\ 19/148 \\ 95/148 \\ 19/148 \end{bmatrix}$$

Issues of Computation of PageRank

- The transition matrix of the Web M is very sparse
- How to use MapReduce for parallel computing?

Change Matrix Representation

$$M = \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$



Source	Degree	Destinations
<i>A</i>	3	<i>B, C, D</i>
<i>B</i>	2	<i>A, D</i>
<i>C</i>	1	<i>A</i>
<i>D</i>	2	<i>B, C</i>

PageRank Iteration Using MapReduce

- Partitioning the matrix M into k^2 blocks, while the vectors are still partitioned into k stripes.
- Need k^2 Map tasks. Each task gets one square of the matrix M , say M_{ij} , and one stripe of the vector v , v_j .
- v is transmitted over the network k times. However, each piece of the matrix is sent only once.

$$\begin{array}{|c|} \hline \mathbf{v}'_1 \\ \hline \mathbf{v}'_2 \\ \hline \mathbf{v}'_3 \\ \hline \mathbf{v}'_4 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline M_{11} & M_{12} & M_{13} & M_{14} \\ \hline M_{21} & M_{22} & M_{23} & M_{24} \\ \hline M_{31} & M_{32} & M_{33} & M_{34} \\ \hline M_{41} & M_{42} & M_{43} & M_{44} \\ \hline \end{array} \begin{array}{|c|} \hline \mathbf{v}_1 \\ \hline \mathbf{v}_2 \\ \hline \mathbf{v}_3 \\ \hline \mathbf{v}_4 \\ \hline \end{array}$$

Topic-Sensitive PageRank

Motivation

- PageRank scores are independent of the queries
- Can we bias PageRank rankings to take into account query keywords?



Topic-sensitive PageRank


To create one vector for each of some small number of topics, biasing the PageRank to favor pages of that topic

Biased Random Walks

- Suppose S is the teleport set
 - a set of integers consisting of the row/column numbers for the pages we have identified as belonging to a certain topic
- Let \mathbf{e}_S be a vector that has 1 in the components in S and 0 in other components.

$$\mathbf{v}' = \beta M \mathbf{v} + (1 - \beta) \mathbf{e}_S / |S|$$

Example



$$\beta M = \begin{bmatrix} 0 & 2/5 & 4/5 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{bmatrix}$$

$\beta = 0.8$

$$\mathbf{v}' = \begin{bmatrix} 0 & 2/5 & 4/5 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{bmatrix} \mathbf{v} + \begin{bmatrix} 0 \\ 1/10 \\ 0 \\ 1/10 \end{bmatrix} \quad \mathcal{S} = \{B, D\}$$

$$\begin{bmatrix} 0/2 \\ 1/2 \\ 0/2 \\ 1/2 \end{bmatrix} \begin{bmatrix} 2/10 \\ 3/10 \\ 2/10 \\ 3/10 \end{bmatrix} \begin{bmatrix} 42/150 \\ 41/150 \\ 26/150 \\ 41/150 \end{bmatrix} \begin{bmatrix} 62/250 \\ 71/250 \\ 46/250 \\ 71/250 \end{bmatrix} \cdots \begin{bmatrix} 54/210 \\ 59/210 \\ 38/210 \\ 59/210 \end{bmatrix}$$

Using Topic-Sensitive PageRank

- To integrate topic-sensitive PageRank into a search engine
 - **Decide** on the **topics** for which we shall create specialized PageRank vectors.
 - **Pick a teleport set for each of these topics**, and use that set to compute the topic-sensitive PageRank vector for that topic.
 - **Find a way of determining the topic or set of topics** that are most **relevant for a particular search query**.
 - **Use the PageRank vectors for that topic or topics** in the ordering of the **responses to the search query**.

The Largest Matrix Computation in the World

- Computing PageRank can be done via matrix multiplication, where the matrix has 3 billion rows and columns.
- The matrix is sparse as average number of outlinks is between 7 and 8.
- Setting $\beta = 0.15$ or above requires at most 100 iterations to convergence.
- Researchers still trying to speed-up the computation.

Monte Carlo Methods in Computing PageRank

- Rather than following a single long random walk, the random surfer can follow many sampled random walks.
- Each walk starts at a random page and either teleports with probability T or continues choosing a link with uniform probability.
- The PR of a page is the proportion of times a sample random walk ended at that page.
- Rather than starting at a random page, start each walk a fixed number of times from each page.

Inlinks and Outlinks

- The number of incoming links to a web page is correlated to the PageRank, but ... this measure is a noisy metric!
- PageRank is biased against new pages (*Googlearchy*), why?
- The long tail of queries, gives less popular web pages some visibility.
- A single outlink to one in the community minimises the loss of PageRank within a community.

Discussion

- Inferring topics from words have been extensively studied for decades
- We can build topic classifiers in advance to guess the topics that are interested

Link Spam

Motivation


- PageRank solves the term spam problem
- Spammers turned to use the **boosting techniques** to fool the PageRank algorithm into overvaluing certain pages

Web Spam Taxonomy


- Defined by Gyongyi and Garcia-Molina [2004]
- Boosting techniques
 - Techniques for achieving high relevance/importance for a web page
- Hiding techniques
 - Techniques to hide the use of boosting
 - From humans and web crawlers

Link Spam by Boosting

- Creating link structures that boost page rank or hubs and authorities scores
- Three kinds of web pages from a spammer's point of view
 - **Inaccessible pages**
 - **Accessible pages**
 - e.g., web log comments pages
 - spammer can post links to his pages
 - **Own pages**
 - Completely controlled by spammer
 - May span multiple domain names



Spam Farm Analysis



Suppose rank contributed by accessible pages = x

Let page rank of target page = y


Rank of each “farm” page = $\beta y / m + (1 - \beta) / n$

$$y = x + \beta m [\beta y / m + (1 - \beta) / n] + (1 - \beta) / n$$

$$= x + \beta^2 y + \beta(1 - \beta)m / n + (1 - \beta) / n$$

$$y = x / (1 - \beta^2) + cm / n \quad \text{where } c = \beta / (1 + \beta)$$

Example



- $y = x/(1-\beta^2) + cm/n$ where $c = \beta/(1+\beta)$
- For $\beta = 0.85$, $1/(1-\beta^2) = 3.6$
 - Multiplier effect for “acquired” page rank
 - By making m large, we can make y as large as we want

Hiding techniques

- Content hiding
 - Use same color for text and page background
- Cloaking
 - Return different pages to crawlers and browsers
- Redirection
 - Alternative to cloaking
 - Redirects are followed by browsers but not crawlers

Combating Link Spam

- Open research area
- Two approaches
 - **TrustRank**
 - a variation of topic-sensitive PageRank designed to lower the score of spam pages.
 - **Spam mass**
 - a calculation that identifies the pages that are likely to be spam and allows the search engine to eliminate those pages or to lower their PageRank strongly.


TrustRank idea

- Basic principle: approximate isolation
 - It is rare for a “good” page to point to a “bad” (spam) page
- Sample a set of “seed pages” from the web
- Have an oracle (human) identify the good pages and the spam pages in the seed set
 - Expensive task, so must make seed set as small as possible

Trust Propagation

- Call the subset of seed pages that are identified as “good” the “trusted pages”
- Set trust of each trusted page to 1
- Propagate trust through links
 - Each page gets a trust value between 0 and 1
 - Use a threshold value and mark all pages below the trust threshold as spam

Example



Rules for trust propagation

- Trust attenuation
 - The degree of trust conferred by a trusted page decreases with distance
- Trust splitting
 - The larger the number of outlinks from a page, the less scrutiny the page author gives each outlink
 - Trust is “split” across outlinks

Simple model

- Suppose trust of page p is $t(p)$
 - Set of outlinks $O(p)$
- For each q in $O(p)$, p confers the trust
 - $\beta t(p) / |O(p)|$ for $0 < \beta < 1$
- Trust is additive
 - Trust of p is the sum of the trust conferred on p by all its inlinked pages
- Note similarity to Topic-Specific Page Rank
 - Within a scaling factor, trust rank = biased page rank with trusted pages as teleport set

Picking the seed set

- Two conflicting considerations
 - Human has to inspect each seed page, so seed set must be as small as possible
 - Must ensure every “good page” gets adequate trust rank, so need make all good pages reachable from seed set by short paths

Approaches to picking seed set

- Suppose we want to pick a seed set of k pages
- PageRank
 - Pick the top k pages by page rank
 - Assume high page rank pages are close to other highly ranked pages
 - We care more about high page rank “good” pages

Inverse page rank

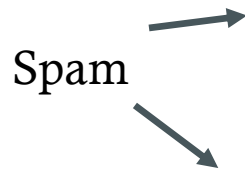
- Pick the pages with the maximum number of outlinks
- Can make it recursive
 - Pick pages that link to pages with many outlinks
- Formalize as “inverse page rank”
 - Construct graph G' by reversing each edge in web graph G
 - Page Rank in G' is inverse page rank in G
- Pick top k pages by inverse page rank

Spam Mass

- To measure each page the fraction of its PageRank that comes from spam
- Suppose page p has PageRank r and TrustRank t
 - The spam mass of p is $(r-t)/r$
- Example

	Node	PageRank	TrustRank	Spam Mass
	<i>A</i>	3/9	54/210	0.229
	<i>B</i>	2/9	59/210	-0.264
	<i>C</i>	2/9	38/210	0.186
	<i>D</i>	2/9	59/210	-0.264

Spam




Hubs and Authorities

HITS (hyperlink-
induced topic search)

Authority vs. Hub

- A is an **authority**
- A is a **hub**




The Intuition Behind HITS

- HITS views important pages as having two flavors of importance
- **Authority pages** are valuable because they provide information about a topic.
- **Hub pages** are valuable because they tell you where to go to find out about that topic.

Pre-processing for HITS

- 1) Collect the top t pages (say $t = 200$) based on the input query; call this the **root set**.
- 2) Extend the root set into a **base set** as follows, for all pages p in the root set:
 - 1) add to the root set all pages that p points to, and
 - 2) add to the root set up-to q pages that point to p (say $q = 50$).
- 3) Delete all links within the same web site in the base set resulting in a **focused sub-graph**.

Expanding the Root Set




Formalizing Hubbiness and Authority

- Vectors \mathbf{h} and \mathbf{a} measure the hubbiness and authority of pages, respectively
- Iterate until Convergence

$$\begin{aligned}\mathbf{h} &= \lambda L \mathbf{a} \\ \mathbf{a} &= \mu L^T \mathbf{h}\end{aligned}\quad L_{ij} = \begin{cases} 1 & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}$$

- λ and μ are unknown constant representing the scaling factors needed
- L is the link matrix of the Web E

Example



$$L = \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$L^T = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

Example

First two iterations

$$\begin{array}{ccccc}
 \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} & \begin{bmatrix} 1 \\ 2 \\ 2 \\ 2 \\ 1 \end{bmatrix} & \begin{bmatrix} 1/2 \\ 1 \\ 1 \\ 1 \\ 1/2 \end{bmatrix} & \begin{bmatrix} 3 \\ 3/2 \\ 1/2 \\ 2 \\ 0 \end{bmatrix} & \begin{bmatrix} 1 \\ 1/2 \\ 1/6 \\ 2/3 \\ 0 \end{bmatrix} \\
 \mathbf{h} & L^T \mathbf{h} & \mathbf{a} & L \mathbf{a} & \mathbf{h}
 \end{array}$$

$$\begin{array}{cccc}
 \begin{bmatrix} 1/2 \\ 5/3 \\ 5/3 \\ 3/2 \\ 1/6 \end{bmatrix} & \begin{bmatrix} 3/10 \\ 1 \\ 1 \\ 9/10 \\ 1/10 \end{bmatrix} & \begin{bmatrix} 29/10 \\ 6/5 \\ 1/10 \\ 2 \\ 0 \end{bmatrix} & \begin{bmatrix} 1 \\ 12/29 \\ 1/29 \\ 20/29 \\ 0 \end{bmatrix} \\
 L^T \mathbf{h} & \mathbf{a} & L \mathbf{a} & \mathbf{h}
 \end{array}$$

Example

The limits are:

$$\mathbf{h} = \begin{bmatrix} 1 \\ 0.3583 \\ 0 \\ 0.7165 \\ 0 \end{bmatrix}$$

$$\mathbf{a} = \begin{bmatrix} 0.2087 \\ 1 \\ 1 \\ 0.7913 \\ 0 \end{bmatrix}$$

Computation of HITS

- The computation of authority scores and hub scores is the same as the computation of the PageRank scores, using **power iteration**.
- If we use \mathbf{a}_k and \mathbf{h}_k to denote authority and hub vectors at the k th iteration, the iterations for generating the final solutions are

$$\mathbf{a}_k = L^T L \mathbf{a}_{k-1}, \mathbf{a}_0 = [1, 1, \dots, 1]$$

$$\mathbf{h}_k = L^T L \mathbf{h}_{k-1}, \mathbf{h}_0 = [1, 1, \dots, 1]$$

The HITS algorithm

HITS-Iterate(G)

$\mathbf{a}_0 = \mathbf{h}_0 = (1, 1, \dots, 1);$

$k = 1$

Repeat

$\mathbf{a}_k = \mathbf{L}^T \mathbf{L} \mathbf{a}_{k-1};$

$\mathbf{h}_k = \mathbf{L} \mathbf{L}^T \mathbf{h}_{k-1};$

normalize \mathbf{a}_k ;

normalize \mathbf{h}_k ;

$k = k + 1;$

until \mathbf{a}_k and \mathbf{h}_k do not change significantly;

return \mathbf{a}_k and \mathbf{h}_k

Strengths and weaknesses of HITS

- **Strength:** its ability to rank pages according to the query topic, which may be able to provide more relevant authority and hub pages.
- **Weaknesses:**
 - **It is easily spammed.** It is in fact quite easy to influence HITS since adding out-links in one's own page is so easy.
 - **Topic drift.** Many pages in the expanded set may not be on topic.
 - **Inefficiency at query time:** The query time evaluation is slow. Collecting the root set, expanding it and performing eigenvector computation are all expensive operations

Applications of HITS

- Search engine querying (speed an issue)
- Finding web communities.
- Finding related pages.
- Populating categories in web directories.
- Citation analysis.

Summary

- In this chapter, we introduced
 - Term Spam
 - The Google Solution to Term Spam
 - PageRank
 - Transition Matrix of the Web
 - Computing PageRank on Strongly Connected Web Graphs
 - The Random Surfer Model
 - Dead Ends

Summary

- In this chapter, we introduced
 - Spider Traps
 - Taxation Schemes
 - Taxation and Random Surfers
 - Efficient Representation of Transition Matrices
 - Topic-Sensitive PageRank
 - Link Spam
 - Hubs and Authorities
 - Recursive Formulation of the HITS Algorithm

Homework 6:Page Rank Implementation with Spark

報告要準時交

Why PageRank?

- Good example of a more complex algorithm
 - Multiple stages of map & reduce
- Benefits from Spark's in-memory caching
 - Multiple iterations over the same data

Basic Idea

- Give pages ranks (scores) based on links to them
 - Links from many pages → high rank
 - Link from a high-rank page → high rank





Image: en.wikipedia.org/wiki/File:PageRank-hi-res-2.png


Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$




Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$




Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$




Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Algorithm


1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Algorithm

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$

Final state:




Scala Implementation

```
val links = // RDD of (url, neighbors) pairs
var ranks = // RDD of (url, rank) pairs

for (i <- 1 to ITERATIONS) {
  val contribs = links.join(ranks).flatMap {
    case (url, (links, rank)) =>
      links.map(dest => (dest, rank/links.size))
  }
  ranks = contribs.reduceByKey(_ + _)
                      .mapValues(0.15 + 0.85 * _)
}

ranks.saveAsTextFile(...)
```

PageRank Performance



Any Questions?